Phase Quantisation from Optical Lattices to the Zeta Landscape

A Unified Framework Connecting Lattice Cryptography, Fresnel Optics, and Analytic Number Theory

Kim Young-Hu

Independent Researcher

April 14, 2026

Abstract

We identify a shared mathematical grammar—*phase quantisation on algebraic structures*—that connects three seemingly disparate domains: CRYSTALS-Kyber lattice-based post-quantum cryptography, free-space Fresnel diffraction, and the phase structure of the Riemann ξ-function on the critical line. The concrete nucleus is QROP-Net, a fully differentiable pipeline that encodes Kyber ciphertext as optical phase patterns and recovers them via deep-learning phase retrieval. Its $\sin^2$-based quantisation loss, which forces continuous phase predictions onto discrete cyclotomic lattice levels, is shown to be structurally isomorphic to the Bohr–Sommerfeld phase quantisation condition $\Delta\nu = m\pi$ that characterises "resonant primes" in the Resonant Prime Law framework. We further demonstrate that the LWE error threshold $\lfloor q/4 \rfloor = 1920$ in Kyber and the σ-lift deviation $\sigma^* - \frac{1}{2}$ from the critical line share the same role as hard error boundaries beyond which signal recovery fails. A multi-group (Polyadic) ODE-based zero detector, which self-calibrates without external ground truth, mirrors QROP-Net's blind test-time optimisation. Seven structural correspondences are formalised, and their mathematical, physical, and computational implications are discussed with explicit acknowledgement of the boundary between analogy and proof.

Note added in revision (2026-04-08). An attempted empirical realisation of the cyclotomic PQO programme on the critical line, using a discriminant network with $L = 4$ cosine PQO and the synthetic ideal phase $\Psi(n)$ of Equation (43) as supervisory target, returns a decisive null: across three seeds the resulting phase distribution fails to discriminate Riemann zeros from off-zero ordinates (χ^2 p-values 0.13/0.53/0.41; sectors $k = 2, 3$ exactly empty). The diagnosis (Remarks 3.2 and 4.2) is that the synthetic Ψ depends only on $\|X\|$ and contains no ξ-zero information whatsoever, so no term of the four-term loss ever injects a t-local zero label into training; Equation (83) (training potential) and Equation (32) (measurement) are operationally disjoint. The text below preserves the synthetic Ψ as the historical object and the cyclotomic PQO as the algebraic target; both are now flagged as awaiting an experiment with a supervised replacement $\Psi(t) := \arg \xi(\frac{1}{2} + it)$ (Equation (44)).

Keywords: Phase quantisation, cyclotomic polynomial ring, CRYSTALS-Kyber, fractional Fourier transform, Riemann ξ-function, Bohr–Sommerfeld condition, resonant primes, σ-lift, Berry phase, Dirichlet L-functions

Contents

1 Introduction **5**

I The Concrete System: QROP-Net **7**

2 QROP-Net: System and Mathematics **7**
2.1 Polynomial Ring and NTT . . . 7
2.2 Phase Modulation: From Integers to Optics . . . 8
2.3 Differentiable Optical Channel: FrFT . . . 8
2.4 Hybrid Loss Orchestration . . . 9
2.4.1 Data Fidelity Loss . . . 9
2.4.2 Periodic Soft Quantisation Loss . . . 9
2.4.3 L^∞ Exponential Barrier Loss . . . 9
2.4.4 Total Variation Loss . . . 9
2.5 Receiver Neural Network and Decoding . . . 9
2.6 Two-Stage Training and Physical-Layer Security . . . 10
2.6.1 Stage 1: Amortised Pre-Training . . . 10
2.6.2 Stage 2: Blind Test-Time Optimisation (TTO) . . . 10

3 The Phase Quantisation Principle **10**

II The Analytic Extension: Resonance Theory **13**

4 Phase Structure of the Completed Zeta Function **13**
4.1 The Riemann ξ-Function . . . 13
4.2 Amplitude–Phase Decomposition on the Critical Line . . . 14
4.3 The Ideal Phase Function $\Psi(n)$ . . . 14
4.4 Connection to the Riemann–Siegel ϑ Function . . . 17

5 Gate Structure: Phase Resonance as Prime Detection **18**
5.1 Gate A: Phase Quantisation (Alignment) . . . 18
5.2 Gate B: Phase Stationarity (Curvature) . . . 18
5.3 Combined Discriminant: Resonant Points . . . 19

6 The σ-Lift and Error Tolerance **20**
6.1 Deviation from the Critical Line . . . 20
6.2 Scattering Interpretation and Unitarity . . . 20
6.3 Berry Phase and Holonomy . . . 21

7 Multi-Group Zero Detection: The Polyadic Frame **21**
7.1 Target Functions and Evaluation Points . . . 21
7.2 Directional Channels . . . 22
7.3 Central Difference Log-Derivative . . . 22
7.4 Gauge ODE for Phase Correction . . . 23
7.5 Final Discriminant Function . . . 23
7.6 Numerical Verification and Convergence . . . 23
7.7 Cross-Validation: Three Independent Axes . . . 24
7.7.1 Axis A: Partial Euler Product Matching . . . 25
7.7.2 Axis B: Explicit Chebyshev $\psi(x)$ Recovery . . . 25
7.7.3 Axis C: Arithmetic Progression Bias Signal . . . 25

8 ■ The F_2-Discriminant as a Hardy-Z Sign-Change Detector **25**
8.1 Hardy Z and the Analytic-Signal Identity . . . 25
8.2 Empirical Validation: 10/10 Zero Detection on $t \in [10, 50]$ . . . 26
8.3 Zero Difficulty Hierarchy and the Lehmer Phenomenon . . . 27

9 ■ Extension to High Zeros: $t \in [100, 200]$ **27**

10 ■ The $|F_2|$ residual: a tentative mesoscopic observable **28**
10.1 Original observations (April 2026) . . . 28
10.2 2026-04-08 partial reinterpretation . . . 29
10.3 2026-04-10 follow-up: the "hard cluster" was a basis-bandwidth artifact . . . 29
10.4 What remains open . . . 33

III Failed Directions and Their Diagnoses **33**

11 ■ Cyclotomic $\cos^2$ PQO training ($L = 2, 4$): Ψ-decoupling and the isolation cascade **34**

12 ■ Phase-Gauge Gradient Descent (PGGD) **35**

IV Synthesis **36**

13 Structural Correspondence Table **36**
13.1 Mathematical Formalisation of the Correspondences . . . 36
13.2 ■ Dynamical and spectral correspondences . . . 37
13.2.1 Maslov correction as the metaplectic sheet shift . . . 37
13.2.2 Gauge ODE critical transition as Kuramoto synchronisation . . . 37
13.2.3 Lehmer phenomenon and GUE spacing statistics . . . 39
13.2.4 Weyl equidistribution and the PGGD angular walk . . . 40

14 ■ Midpoint Curvature Nonlocality **40**

15 Discussion **42**
15.1 What the Correspondence Means 42
15.2 What the Correspondence Does Not Mean 42
15.3 Predictive Power and Future Directions 43

16 Related Work **47**

17 Conclusion **48**

A QROP-Net Full Equation Reference **50**

B Numerical Tables for Resonance Gates **50**

C Generalisation to Dirichlet L-Functions **50**
C.1 Numerical Verification of Bundle Properties 51
C.2 Blind Zero Prediction 52

D Negative Result: Gram Point Curvature **52**

E Notation Index **53**

1 Introduction

The polynomial ring $R_q = \mathbb{Z}_q[X]/(X^n + 1)$ with $n = 256$ and $q = 7681$ is the algebraic backbone of CRYSTALS-Kyber, the NIST-standardised post-quantum key encapsulation mechanism [1]. The defining relation $X^n + 1 = 0$—or equivalently $X^n \equiv -1$, the origin of our project codename $k^n = -1$—is a cyclotomic polynomial whose roots are the $2n$-th roots of unity:

$$\omega_j = e^{i\pi(2j+1)/n}, \quad j = 0, 1, \ldots, n-1. \tag{1}$$

These n phases are equally spaced on the unit circle, separated by $\Delta\theta = \pi/n$. This discrete, periodic phase structure is not merely an algebraic convenience for fast polynomial multiplication via the Number Theoretic Transform (NTT); it is, we argue, an instance of a more general *phase quantisation principle* that appears in at least two other domains:

1. **Fresnel optics**: The fractional Fourier transform (FrFT) maps a coherent wavefront $U_0 = e^{i\phi}$ through free-space diffraction, where the fractional order α encodes the propagation distance as a phase parameter.

2. **Analytic number theory**: The completed Riemann ξ-function $\xi(s) = \frac{1}{2}s(s-1)\pi^{-s/2}\Gamma(s/2)\zeta(s)$ decomposes on the critical line $s = \frac{1}{2} + it$ into amplitude $A(t) = |\xi(\frac{1}{2}+it)|$ and phase $\Phi_\xi(t) = \arg \xi(\frac{1}{2}+it)$, with phase quantisation conditions intimately linked to the distribution of primes and the Riemann Hypothesis.

Thesis. The $\sin^2$-based quantisation loss of QROP-Net, the Bohr–Sommerfeld condition $\Delta\nu = m\pi$ on the critical line, and the cyclotomic root structure $X^n \equiv -1$ are three manifestations of a single mathematical grammar: *periodic phase penalty on an algebraic structure forces continuous signals onto discrete lattice levels, with a hard error boundary beyond which decoding fails.*

Contributions.

1. We present QROP-Net, a fully differentiable pipeline unifying Kyber lattice cryptography, Fresnel optics, and deep phase retrieval (Part I).

2. We abstract the phase quantisation mechanism into a general principle bridging discrete and continuous Fourier analysis (Section 3).

3. We apply this principle to the phase structure of $\xi(s)$, formulating a two-gate resonance criterion for prime detection (Part II, Sections 4–5).

4. We establish a structural correspondence between the LWE error threshold and the σ-lift from the critical line (Section 6).

5. We present a multi-group ODE zero detector that self-calibrates without ground truth, mirroring QROP-Net's blind test-time optimisation (Section 7).

6. We formalise seven structural correspondences and rigorously delimit the boundary between analogy and proof (Part III, Section 13).

7. **(Negative empirical result, 2026-04-08)** We attempt the cyclotomic PQO programme empirically with $L = 4$ on $t \in [100, 200]$, document a complete failure to discriminate zeros from non-zeros across three seeds, trace the cause to a Ψ-*decoupling defect* of Equation (43), and propose the supervised replacement Equation (44) as the necessary precondition for any future positive result (Part III, Remarks 3.2 and 4.2).

8. **(Isolation experiment, 2026-04-09)** We isolate the target-injection, loss-weighting, and read-out channels and demonstrate that, with all three closed, the supervised chain $\arg \xi(\frac{1}{2} + it) \to \mathcal{L}_{\text{tgt}} \to \hat{\Psi}$ is sound at off-zero sampling points ($p = 1.18 \times 10^{-15}$, ratio $18.2\times$), localising the remaining defect to the network's output representation near zeros (Remark 4.3).

Reading Guide: Evidence Tiers

This work is a *consolidated research notebook*: it records the full trajectory of a four-month investigation including ideas that remain conjectural and empirical attempts that failed. We therefore classify every substantive claim into one of three tiers, signalled explicitly in each section heading and in the table of contents:

- **Tier 1 — Validated.** Claims backed by complete derivations and successful numerical experiments. Examples: the F_2-discriminant zero detector (Section 8), the supervised identification $\Psi(t) = \arg \xi(\frac{1}{2} + it)$ via the isolation experiment (Remark 4.3), the high-zero extension to $t \in [100, 200]$ (Section 9). These sections should be read as mature results.

- **Tier 2 — Ideas and partial results.** Analytic correspondences, theoretical constructions, or experiments with suggestive but incomplete evidence. Examples: Gate B stationarity (Section 5), the σ-lift (Section 6), the seven structural correspondences (Section 13), the $|F_2|$ spatial observables with their 2026-04-08 reinterpretation (Section 10). These sections motivate and sketch; they should be read as research programme rather than settled theory.

- **Tier 3 — Failed attempts and their diagnoses.** Empirical programmes that did not succeed, preserved here with full diagnostic analysis. Examples: cyclotomic $\cos^2$ $L \in \{2, 4\}$ PQO training (Section 11), the phase-gauge gradient descent (PGGD) optimiser (Section 12). We retain these sections, compressed but intact, because the diagnostic chain is itself a contribution: it documents exactly which combinations of target design, loss weighting, measurement, and model expressivity are incompatible, saving future work from the same failure modes. A failed direction whose cause is understood is a partial positive result.

We will often refer to the tier of a given section or remark by the symbols ■/ ■/■. The reader may use the tier as a calibration for how much trust to place in a given statement. The structural division of the paper into Parts I–IV respects this tiering: Part I and Part II collect Tier 1 and Tier 2 material, Part III collects Tier 3 failure diagnoses, and Part IV integrates all three into a unified picture.

Part I

The Concrete System: QROP-Net

2 QROP-Net: System and Mathematics

2.1 Polynomial Ring and NTT

Kyber operates over $R_q = \mathbb{Z}_q[X]/(X^n+1)$ where $q = 7681$ is chosen to satisfy $q \equiv 1 \pmod{2n}$, enabling the Number Theoretic Transform (NTT) for $O(n \log n)$ polynomial multiplication. The ring condition $X^n + 1 = 0$ factors as:

$$X^n + 1 = \prod_{j=0}^{n-1} (X - \omega_j), \quad \omega_j = e^{i\pi(2j+1)/n} \in \mathbb{C}. \tag{2}$$

Every element $a(X) \in R_q$ corresponds to a vector $\mathbf{a} = (a_0, \ldots, a_{n-1}) \in \mathbb{Z}_q^n$, and multiplication $a \cdot b \pmod{X^n + 1}$ becomes component-wise multiplication in the NTT domain.

Definition 2.1 (Kyber IND-CPA Encryption)**.** Given public key $(\mathbf{A}, \mathbf{t})$ where $\mathbf{A} \in R_q^{k \times k}$, $\mathbf{t} \in R_q^k$, and message $m \in \{0,1\}^n$, the ciphertext $(\mathbf{u}, v)$ is:

$$\mathbf{u} = \mathbf{A}^\top \mathbf{r} + \mathbf{e}_1, \tag{3}$$

$$v = \mathbf{t}^\top \mathbf{r} + e_2 + \lceil q/2 \rfloor \cdot m, \tag{4}$$

where $\mathbf{r}, \mathbf{e}_1, e_2 \sim \beta_\eta$ are sampled from the centred binomial distribution:

$$\beta_\eta : \quad x = \sum_{i=1}^{\eta} (a_i - b_i), \quad a_i, b_i \sim \text{Bernoulli}(\tfrac{1}{2}), \quad x \in [-\eta, \eta]. \tag{5}$$

Definition 2.2 (Compression and Decompression)**.**

$$\text{Compress}_q(x, d) = \left\lceil \frac{2^d}{q} \cdot x \right\rfloor \bmod 2^d, \tag{6}$$

$$\text{Decompress}_q(x, d) = \left\lceil \frac{q}{2^d} \cdot x \right\rfloor, \tag{7}$$

where $\lceil \cdot \rfloor$ denotes ties-rounded-up: $\lfloor x + 0.5 \rfloor$. Parameters: $d_u = 11$ (2048 levels for $\mathbf{u}$), $d_v = 3$ (8 levels for v).

Theorem 2.1 (LWE Correctness Criterion)**.** *Decryption succeeds if and only if the total error $w = v - \mathbf{s}^\top \mathbf{u} \pmod{q}$ satisfies:*

$$\|w - \lceil q/2 \rfloor \, \hat{m}\|_\infty < \lceil q/4 \rfloor = 1920. \tag{8}$$

A single coefficient exceeding this bound causes complete decryption failure ($\text{DFR} = 1$*). The ring distance is:*

$$|e|_{\mathbb{Z}_q} = \min(e,\ q - e), \quad e \in \{0, 1, \ldots, q-1\}. \tag{9}$$

2.2 Phase Modulation: From Integers to Optics

Each compressed ciphertext coefficient $c_i \in \{0, \ldots, 2^{d_u} - 1\}$ is mapped to a phase angle and encoded as a unit-amplitude complex wavefront:

$$\phi_i = \frac{2\pi \cdot c_i}{2^{d_u}}, \qquad U_0(x, y) = e^{i\phi(x,y)}, \quad |U_0| \equiv 1. \tag{10}$$

The $2^{d_u} = 2048$ discrete ciphertext levels map to 2048 equally spaced phases on the unit circle, separated by $\Delta\phi = 2\pi/2048$. This is a direct physical realisation of the algebraic structure: the polynomial ring's discrete coefficients become discrete optical phases on the SLM.

A dimensional constraint links the cryptographic and optical domains:

$$N = \sqrt{k \cdot n + n} \times M = \sqrt{1024} \times 8 = 256, \tag{11}$$

where $M = 8$ is the macro-pixel expansion factor for noise resilience.

2.3 Differentiable Optical Channel: FrFT

When the wavefront U_0 propagates a distance d through free space, the diffraction integral is equivalent to a fractional Fourier transform of order p:

$$p = \frac{2}{\pi} \arctan\left(\frac{\lambda d}{N \Delta x^2}\right), \qquad \alpha = \frac{p\pi}{2}, \tag{12}$$

where $\lambda = 500\,\text{nm}$ is the wavelength and $\Delta x = 8\,\mu\text{m}$ is the pixel pitch. The propagation distance d is declared as a learnable parameter, enabling gradient-based auto-focusing:

$$d \xrightarrow{\partial} R = \frac{\lambda d}{N \Delta x^2} \;\rightarrow\; \alpha = \arctan(R) \;\rightarrow\; s_2 = \sqrt{1 + R^2}. \tag{13}$$

The IP-DFrFT is implemented via chirp–FFT–chirp decomposition, reducing complexity from $O(N^4)$ to $O(N^2 \log N)$:

$$\boxed{U_d = \underbrace{(1 - i\cot\alpha)}_{A_\alpha^2} \cdot \mathcal{F}^{-1}\Big[\mathcal{F}\big[\underbrace{U_0 \cdot e^{i\pi(\cot\alpha - \csc\alpha)\, r^2}}_{\text{input chirp}}\big] \;\cdot\; \mathcal{F}\big[\underbrace{e^{i\pi\, \csc\alpha\, r'^2}}_{\text{kernel chirp}}\big]\Big] \cdot \Delta x^2} \tag{14}$$

where $r^2 = (m^2 + n^2)\Delta x^2$ on $N \times N$ and r'^2 on zero-padded $2N \times 2N$ to prevent circular aliasing.

The sensor captures intensity and adds physical noise:

$$I_{\text{clean}} = \frac{|U_d|^2}{1 + \cot^2\alpha}, \tag{15}$$

$$I_{\text{shot}} \sim \text{Poisson}(I_{\text{clean}} \cdot \text{FWC}), \tag{16}$$

$$I_{\text{meas}} = \text{ADC}_{\text{12-bit}}(I_{\text{shot}} + \mathcal{N}(0, \sigma_{\text{read}}^2)). \tag{17}$$

The straight-through estimator (STE) preserves gradient flow through the non-differentiable

Poisson sampling and ADC rounding:

$$\mathrm{STE}(x_{\text{noisy}}, x_{\text{clean}}) = x_{\text{clean}} + \mathrm{stopgrad}(x_{\text{noisy}} - x_{\text{clean}}). \tag{18}$$

2.4 Hybrid Loss Orchestration

Four loss terms are combined under a curriculum scheduler.

2.4.1 Data Fidelity Loss

The predicted wavefront $\hat{U}_0$ is propagated through the FrFT to produce $\hat{I}$. Comparison in the amplitude domain ensures uniform gradient flow:

$$\mathcal{L}_{\text{fid}} = \frac{1}{N^2} \sum_{x,y} \left(\sqrt{\hat{I}(x,y) + \epsilon} - \sqrt{I_{\text{target}}(x,y) + \epsilon} \right)^2. \tag{19}$$

2.4.2 Periodic Soft Quantisation Loss

A $\sin^2$-based penalty forces continuous phases onto discrete Kyber lattice levels:

$$\boxed{\mathcal{L}_{\text{quant}} = \frac{1}{2} \left[\frac{1}{kn} \sum_{i=1}^{kn} \sin^2\!\left(\frac{2^{d_u}}{2} \phi_i^{(u)} \right) + \frac{1}{n} \sum_{j=1}^{n} \sin^2\!\left(\frac{2^{d_v}}{2} \phi_j^{(v)} \right) \right]} \tag{20}$$

The $\sin^2$ function creates a *valley–ridge landscape*: zero loss at exact lattice positions $\phi = 2\pi k/2^d$, maximum penalty at midpoints. This is the key equation that bridges to analytic number theory (Section 3).

2.4.3 L^∞ Exponential Barrier Loss

An exponential barrier enforces the LWE threshold on the most dangerous outliers:

$$\mathcal{L}_{\max} = \frac{1}{|S_K|} \sum_{i \in S_K} \exp(\gamma \cdot \mathrm{ReLU}(|e_i|_{\mathbb{Z}_q} - \tau_{\text{safe}})), \tag{21}$$

where $\tau_{\text{safe}} = 0.8 \times 1920 = 1536$, $\gamma = 10$, and S_K is the set of top-5% error pixels.

2.4.4 Total Variation Loss

On the complex wavefront (not the phase angle, to avoid wrapping artefacts):

$$\mathcal{L}_{\text{TV}} = \frac{1}{N^2} \sum_{x,y} (|\hat{U}_0^{x+1,y} - \hat{U}_0^{x,y}|^2 + |\hat{U}_0^{x,y+1} - \hat{U}_0^{x,y}|^2). \tag{22}$$

2.5 Receiver Neural Network and Decoding

A LeWin U-Net (window-based multi-head self-attention + locally-enhanced feed-forward) maps the intensity measurement I_{meas} to a 2-channel $(\mathrm{Re}, \mathrm{Im})$ output, from which the phase

is extracted via:

$$\hat{\phi} = \operatorname{atan2}(\mathrm{Im},\ \mathrm{Re}) \bmod 2\pi, \tag{23}$$

$$\hat{U}_0 = e^{i\hat{\phi}}, \quad |\hat{U}_0| \equiv 1. \tag{24}$$

Macro-pixel pooling uses circular statistics to avoid phase averaging catastrophe:

$$\bar{\phi} = \operatorname{atan2}\left(\frac{1}{M^2}\sum \sin\phi_{ij},\ \frac{1}{M^2}\sum \cos\phi_{ij}\right) \bmod 2\pi. \tag{25}$$

Final Kyber decapsulation:

$$\hat{\mathbf{u}}, \hat{v} = \mathrm{Decompress}_q(c_{\mathrm{pred}}, d_u, d_v), \tag{26}$$

$$w = \hat{v} - \mathbf{s}^\top \hat{\mathbf{u}} \pmod q, \tag{27}$$

$$\hat{m}_j = \left\lceil \frac{2}{q} \cdot w_j \right\rfloor \bmod 2. \tag{28}$$

2.6 Two-Stage Training and Physical-Layer Security

2.6.1 Stage 1: Amortised Pre-Training

Synthetic ciphertexts are generated on-the-fly; AdamW jointly updates 119 U-Net parameters and the physical distance d.

2.6.2 Stage 2: Blind Test-Time Optimisation (TTO)

At deployment, only physics-based and mathematics-based constraints are available:

$$\mathcal{L}_{\mathrm{blind}} = \mathcal{L}_{\mathrm{fid}} + 0.5\,\mathcal{L}_{\mathrm{quant}} + 0.05\,\mathcal{L}_{\mathrm{TV}}. \tag{29}$$

The U-Net backbone is frozen at $\mathrm{lr} = 10^{-5}$ while d receives $\mathrm{lr} = 10^{-3}$, enabling rapid auto-focusing in only 15 gradient steps *without access to the ground-truth ciphertext.*

Theorem 2.2 (Avalanche Effect — Physical-Layer Security)**.** *An eavesdropper who possesses identical U-Net weights, PRNG seed, and Kyber secret key, but estimates* $d' = d + 0.1\,\mathrm{mm}$*, experiences* DFR $= 1$*. The FrFT angle* α *is a nonlinear function of* d*; a 0.1 mm offset cascades through the chirp–FFT–chirp pipeline, causing* $\|e\|_\infty > 1920$ *for every ciphertext element.*

3 The Phase Quantisation Principle

We now abstract the $\sin^2$ quantisation mechanism (Equation (20)) into a general principle.

Definition 3.1 (Phase Quantisation Operator)**.** Let $\phi \in [0, 2\pi)$ be a continuous phase variable and $L \in \mathbb{N}$ a target number of discrete levels. The *phase quantisation penalty* is:

$$\mathcal{P}(\phi; L) = \sin^2\left(\frac{L}{2}\,\phi\right). \tag{30}$$

This function has the following properties:

1. $\mathcal{P}(\phi; L) = 0$ if and only if $\phi \in \{2\pi k/L : k = 0, 1, \ldots, L-1\}$ (the lattice positions).
2. $\mathcal{P}(\phi; L) = 1$ at the midpoints $\phi = 2\pi(k + \frac{1}{2})/L$.
3. $\nabla_\phi \mathcal{P} = \frac{L}{2}\sin(L\phi)$ provides a gradient that steers ϕ toward the nearest lattice point.
4. The landscape is $2\pi/L$-periodic, creating L equally spaced attraction basins on the unit circle.

In QROP-Net, the lattice levels are $L = 2^{d_u} = 2048$ (for $\mathbf{u}$-coefficients) and $L = 2^{d_v} = 8$ (for v-coefficients). The $\sin^2$ loss drives each continuous phase prediction toward the nearest discrete Kyber level, realising a *soft quantisation* that is fully differentiable.

Remark 3.1 (Cyclotomic Connection)**.** The L zeros of $\mathcal{P}(\phi; L)$ are precisely the L-th roots of unity $e^{2\pi i k/L}$. When $L = 2n = 512$ (the order of the cyclotomic polynomial $X^n + 1$), these are the roots $\omega_0, \ldots, \omega_{n-1}$ of Equation (1). The $\sin^2$ loss is thus the *analytic restatement* of the algebraic condition $X^n + 1 = 0$ in the phase domain.

Remark 3.2 (Cosine PQO and Higher-Resolution Cyclotomic Lattices)**.** The PQO admits a structurally identical *cosine* variant

$$\mathcal{P}_{\cos}(\phi; L) = \cos^2\left(\frac{L}{2}\,\phi\right), \tag{31}$$

whose zero set is $\phi \in \{(2k+1)\pi/L : k = 0, 1, \ldots, L-1\}$. These are precisely the *primitive* $2L$*-th roots of unity*, i.e. the L roots ω_j of $X^L + 1 = 0$ from Equation (1). Hence $\mathcal{P}_{\cos}(\,\cdot\,; L)$ is the L-th cyclotomic PQO, parallel to Remark 3.1 but realising the algebraic constraint $X^L + 1 = 0$ *directly* on the phase domain rather than the shifted unity-root constraint $X^L - 1 = 0$ associated with $\sin^2$.

Trivial collapse at $L = 2$. At $L = 2$ the attractor set degenerates to $\{i, -i\}$. Any signal trained with a non-trivially weighted $\mathcal{P}_{\cos}(\phi; 2)$ term collapses globally onto one of these two points: in particular, in a critical-line discriminant network whose loss includes such a term, the empirical phase $\arg Z(t)$ is driven to $\pi/2$ *for all t in the trained band*, irrespective of whether t lies near a Riemann zero. The $L = 2$ cosine PQO therefore carries no zero-localising information; it merely confirms that the attractor mechanism functions as designed.

Higher-resolution generalisation (attempted). For $L \geq 4$ the lattice acquires sufficient cardinality, in principle, to act as a sector classifier on S^1. Each candidate ordinate t is assigned a discrete label

$$k(t) \;=\; \arg\min_{0 \leq k < L} \left|\arg Z(t) - (2k{+}1)\pi/L\right| \pmod{2\pi}, \tag{32}$$

naming the cyclotomic sector $\omega_{k(t)}$ closest to the trained phase, with the hope that Riemann zeros would occupy a non-trivial sub-lattice and thereby provide an arithmetic stratification finer than the $L = 2$ case admits.

Empirical refutation at $L = 4$ (2026-04-08). A direct test of this programme was performed using a critical-line discriminant network with $L = 4$, hidden width 32, three seeds

$\{42, 123, 777\}$, 200 training epochs, and a labelled population of 50 Riemann zeros versus 200 off-zero ordinates of gap ≥ 0.5 (training script `scripts/phase_L4_experiment.py` in the companion GDL_UNIFIED repository; output dump at `outputs/analysis/phase_L4_experiment.{json,txt}`). The result is unambiguously negative on three independent axes:

1. All three seeds fail a χ^2 test of independence between the zero and off-zero k-distributions ($p = 0.127, 0.528, 0.410$; pooled $p = 0.154$).
2. Of the four cyclotomic sectors $\{k = 0, 1, 2, 3\}$ only sectors $k = 0$ and $k = 1$ are ever populated; sectors $k = 2, 3$ contain *exactly zero* samples on every seed (asymmetric basin collapse).
3. Mean nearest-attractor residuals for zeros and off-zeros agree to within $\sim 5\%$ across all seeds.

The L=2 trivial collapse and the L=4 asymmetric collapse therefore share a single root cause; we now identify it.

Diagnosis: Ψ-decoupling defect. The training loss of the discriminant network is a sum $\mathcal{L} = \lambda_{\text{res}}\mathcal{L}_{\text{res}} + \lambda_{\text{curv}}\mathcal{L}_{\text{curv}} + \lambda_{\text{tgt}}\mathcal{L}_{\text{tgt}} + \lambda_{\text{pqo}}\mathcal{L}_{\text{pqo}}$, in which the target-matching term $\mathcal{L}_{\text{tgt}}$ is built against Ψ as defined in Definition 4.3 / Equation (43). As written, Ψ depends only on $\|X\|$ and is a synthetically chosen $(\alpha, \beta, \gamma, \delta)$ curve carrying *no ξ-zero information whatsoever*. Consequently:

- No term of $\mathcal{L}$ injects a t-local zero label into training;
- The cosine PQO term $\mathcal{L}_{\text{pqo}}$ is a *t-independent* potential whose basin selection is fixed by initialisation, not by the location of t relative to a zero;
- The labelling rule Equation (32) is therefore an *operationally disjoint measurement* placed on top of a network that has never been told what a zero is.

In particular Equation (83) (training) and Equation (32) (measurement) act on incompatible signal categories: the former *erases* the t-localised information that the latter is trying to read out. Increasing L from 2 to $4, 8, \ldots$ cannot repair this — the broken link is upstream of L.

Status of the cyclotomic PQO programme. The $L \geq 4$ programme as previously formulated is suspended. Its revival requires *replacing* Equation (43) by a target carrying genuine zero information (the natural candidates are $\Psi(t) := \arg \xi(\frac{1}{2} + it)$ or the Hardy phase $\arg Z(t)$ of Equation (46)), demoting $\mathcal{P}_{\cos}$ from a *training* potential to a pure *measurement* of the resulting phase, and re-running the $L = 4$ protocol under the new supervision. Until such an experiment exists, no claim of cyclotomic stratification of Riemann zeros is supported by this work; the data files `phase_L4_experiment.*` and `phase_pi_half_verify.txt` are retained as deterministic negative controls and are part of the record on which the diagnosis above rests.

Remark 3.3 (Bohr–Sommerfeld Analogy)**.** The condition $\mathcal{P}(\phi; L) = 0$ is equivalent to $L\phi/2 = m\pi$ for $m \in \mathbb{Z}$, *i.e.*,

$$\Delta\nu \;=\; m\pi, \qquad m \in \mathbb{Z}, \tag{33}$$

where $\Delta\nu = L\phi/2$ is a rescaled phase accumulation. This is precisely the form of the *Bohr–Sommerfeld quantisation condition* in semiclassical mechanics, where the action integral around a closed orbit must equal an integer multiple of $\pi\hbar$ (in appropriate units). In optics, the same condition determines constructive interference: the path length must accommodate an integer number of half-wavelengths. The phase quantisation principle thus unifies:

Domain	**Quantised Variable**	**Condition**
QROP-Net	Ciphertext phase ϕ_i	$\sin^2(2^{d_u-1}\phi_i) = 0$
Semiclassical QM	Action integral $\oint p\,dq$	$= (n + \frac{1}{2})\pi\hbar$
Optics (cavity)	Optical path length nL_{cav}	$= m\lambda/2$
Critical line	Phase of $\xi(s)$	$\Delta\nu = m\pi$

Part II

The Analytic Extension: Resonance Theory

Every mathematical claim in this part has either a concrete QROP-Net parallel (noted in Remark environments) or an independent reference in the literature. Claims from earlier versions of the Resonant Prime Law that lack both are omitted.

4 Phase Structure of the Completed Zeta Function

4.1 The Riemann ξ-Function

Definition 4.1 (Completed Zeta Function)**.** The Riemann ξ-function is defined as:

$$\xi(s) = \frac{1}{2}s(s-1)\pi^{-s/2}\Gamma\left(\frac{s}{2}\right)\zeta(s), \tag{34}$$

where $\zeta(s) = \sum_{n=1}^{\infty} n^{-s}$ is the Riemann zeta function and Γ is the Euler gamma function. The function ξ is entire (no poles) and satisfies the functional equation:

$$\xi(s) = \xi(1-s). \tag{35}$$

On the critical line $s = \frac{1}{2} + it$, ξ is real-valued: $\xi(\frac{1}{2} + it) \in \mathbb{R}$.

Definition 4.2 (Weierstrass Product Representation)**.** Since $\xi(s)$ is entire, it admits a Weierstrass product over its zeros ρ (the nontrivial zeros of ζ):

$$\xi(s) = e^{g(s)} \prod_{\rho} \left(1 - \frac{s}{\rho}\right), \tag{36}$$

for some entire function $g(s)$ ensuring convergence. Differentiating $\ln \xi$ gives the *logarithmic derivative*:

$$\frac{\xi'(s)}{\xi(s)} = g'(s) + \sum_{\rho} \frac{1}{s-\rho}. \tag{37}$$

Away from the zeros, $g'(s)$ is smooth, and the dominant contributions come from the poles at $s = \rho$—the nontrivial zeros of ζ.

4.2 Amplitude–Phase Decomposition on the Critical Line

For general $s = \sigma + it$ (not restricted to $\sigma = \frac{1}{2}$), we decompose ξ into amplitude and phase:

$$A(\sigma, t) = |\xi(\sigma + it)|, \tag{38}$$

$$v(\sigma, t) = \arg \xi(\sigma + it), \tag{39}$$

so that $\xi(\sigma + it) = A(\sigma, t)\, e^{iv(\sigma,t)}$. Writing $\log \xi = u + iv$ where $u = \log |\xi|$ and $v = \arg \xi$, the Cauchy–Riemann equations give:

$$v_\sigma = \mathrm{Re}\left(\frac{\xi'}{\xi}\right), \quad v_t = -\mathrm{Im}\left(\frac{\xi'}{\xi}\right), \quad u_t = -v_\sigma, \quad u_\sigma = v_t. \tag{40}$$

On the critical line $\sigma = \frac{1}{2}$, we write:

$$\Phi_\xi(t) \;:=\; v(\tfrac{1}{2}, t) = \arg \xi(\tfrac{1}{2} + it), \tag{41}$$

$$A(t) \;:=\; A(\tfrac{1}{2}, t) = |\xi(\tfrac{1}{2} + it)|. \tag{42}$$

Since $\xi(\frac{1}{2} + it) \in \mathbb{R}$, the phase $\Phi_\xi(t)$ takes values in $\{0, \pi\}$ (modulo 2π), jumping by π at each zero crossing.

Remark 4.1 (QROP-Net Parallel: Amplitude–Phase Separation)**.** In QROP-Net, the transmitted wavefront $U_0 = e^{i\phi}$ has unit amplitude ($|U_0| \equiv 1$) and discrete phase ϕ. The receiver must recover ϕ from the intensity $I = |U_d|^2$, which involves "undoing" the amplitude–phase mixing caused by diffraction. The structure $\xi = A\, e^{iv}$ on the critical line presents the identical mathematical challenge: separating the amplitude envelope from the underlying phase structure.

4.3 The Ideal Phase Function $\Psi(n)$

Reader's warning. The synthetic form Equation (43) below was empirically falsified as a supervisory target on 2026-04-08; see Remark 4.2 immediately following the definition and Remark 3.2 for the diagnosis. The definition is retained for historical exposition; the operational replacement used in subsequent revisions is Equation (44).

Definition 4.3 (Smooth Prime-Counting Phase)**.** For integer $n \geq 2$, the ideal phase function is defined as:

$$\Psi(n) = \alpha \log n + \frac{\beta \sin(\gamma \log n)}{n^\delta}, \tag{43}$$

where $\alpha, \beta, \gamma, \delta$ are structural constants:

- $\alpha > 0$ determines the logarithmic ascent (alignment slope),
- $\beta \geq 0$ controls the oscillation amplitude (resonance strength),

- $\gamma > 0$ is the phase rotation frequency (related to the local Gamma factor),
- $\delta > 0$ controls exponential damping: the oscillation term decays as $n^{-\delta}$, so $\Psi(n) \to \alpha \log n$ for large n.

The monotone term $\alpha \log n$ mirrors the Riemann–von Mangoldt formula $N(T) \sim \frac{T}{2\pi} \log \frac{T}{2\pi e}$ for the zero counting function, while the oscillatory correction $\beta \sin(\gamma \log n)/n^{\delta}$ is intended as a heuristic stand-in for local fluctuations related to prime gaps.

Remark 4.2 (Decoupling defect of Equation (43); documented 2026-04-08)**.** The form Equation (43) as written is a synthetic $(\alpha, \beta, \gamma, \delta)$ curve depending only on the magnitude n (or, in the network implementation, on $\|X\|$); it carries *no* information about the location of Riemann zeros. When used as a target for a discriminant network — as in the L=4 cyclotomic PQO experiment of Remark 3.2 — none of the four loss terms $\mathcal{L}_{\text{res}}, \mathcal{L}_{\text{curv}}, \mathcal{L}_{\text{tgt}}, \mathcal{L}_{\text{pqo}}$ injects a t-local zero signal into training, and the resulting network necessarily fails to discriminate zeros from non-zeros regardless of L. The defect is structural rather than parametric: no choice of $(\alpha, \beta, \gamma, \delta)$ can repair it.

The intended replacement, used in subsequent revisions of this work, is to redefine the ideal phase on the critical line directly,

$$\Psi(t) \;:=\; \arg \xi(\tfrac{1}{2} + it), \tag{44}$$

or equivalently to take $\Psi(t) := \arg Z(t)$ via the Hardy Z-function of Equation (46). Either choice supplies ξ-zero information *by construction*, restoring the causal chain between the supervisory target and the cyclotomic measurement Equation (32). The synthetic form Equation (43) is retained in the text only as the historical object whose failure motivated the diagnosis.

Supervised-Ψ control experiment (2026-04-08). The replacement Equation (44) was tested directly. The same discriminant network architecture (hidden width 32, three layers, three seeds $\{42, 123, 777\}$, 200 epochs, $t \in [100, 200]$, 50 zeros vs. 200 off-zero ordinates of gap ≥ 0.5) was retrained with two single-line modifications relative to the L=4 baseline: (i) the model's Ψ output was overwritten by the precomputed $\arg \xi(\frac{1}{2} + it)$ supplied as a supervisory signal; (ii) $\lambda_{\text{pqo}} = 0$, removing the cosine PQO from training and retaining it only as the cyclotomic measurement Equation (32) (script `scripts/phase_supervised_experiment.py`; output `phase_supervised_experiment.{json,txt}`). The result is again negative on the headline metric — the seed-pooled phase MSE of $\arg Z(t)$ against $\arg \xi(\frac{1}{2} + it)$ is essentially identical for zeros and non-zeros (means 2.256 vs. 2.247, ratio 1.004, Welch $p = 0.976$) — and the per-seed Welch tests are inconsistent in direction (one seed $p = 0.034$ *against* the expected ordering). Two structural symptoms are nevertheless new and informative:

1. The asymmetric basin collapse of the L=4 baseline is broken: all four cyclotomic sectors $k \in \{0, 1, 2, 3\}$ are now populated on every seed. The t-independent attractor freeze induced by $\mathcal{L}_{\text{pqo}}$ is therefore confirmed, by ablation, as one of the two contributors to that collapse.

2. The recorded best validation loss (≈ 1.5–1.7) is substantially smaller than the recorded phase MSE (≈ 2.25), implying that the $\mathcal{L}_{\text{tgt}}$ contribution is a small fraction of $\mathcal{L}$ even after $\mathcal{L}_{\text{pqo}}$ has been zeroed out. The remaining $\mathcal{L}_{\text{res}}$ and $\mathcal{L}_{\text{curv}}$ terms dominate the gradient, so the network is not actually being asked to fit $\arg \xi$.

This refines, but does *not* retract, the diagnosis above: Ψ-decoupling is a necessary defect of Equation (43), and Equation (44) repairs it at the level of the target alone; but the loss-weight imbalance and a suspected measurement artefact (arithmetic averaging of channel-wise $\arg Z$, which destroys the $\pm\pi$ wrap signal that Equation (41) predicts) are *additional* broken links. A second control experiment with $\lambda_{\text{res}} = \lambda_{\text{curv}} = \lambda_{\text{pqo}} = 0$ (only $\mathcal{L}_{\text{tgt}}$ active) and a circular-mean read-out is described in Remark 4.3 below.

Remark 4.3 (Isolation experiment: supervisory chain is sound, $\pm\pi$ jump representability is not; documented 2026-04-09)**.** To disentangle the three suspected defects of Remark 4.2 — (i) Ψ-decoupling via Equation (43), (ii) loss-weight imbalance, and (iii) an arithmetic-mean read-out that cannot represent the $\pm\pi$ wrap — we ran a maximally isolated control in which all three channels were simultaneously closed. The network was retrained with $\lambda_{\text{res}} = \lambda_{\text{curv}} = \lambda_{\text{pqo}} = 0$ and only $\mathcal{L}_{\text{tgt}}$ active, the supervisory target hard-set to $\arg \xi(\frac{1}{2} + it)$ as in Equation (44), and the per-sample phase averaged via a *circular* mean $\arg \sum_c e^{i\hat{\Psi}_c}$ rather than the arithmetic channel-wise average (script `scripts/phase_isolation_experiment.py`; output `phase_isolation_experiment.{json,txt}`; same architecture, seeds, epochs, and t-window as in Remarks 3.2 and 4.2). The result is strongly *positive*:

- **Circular read-out.** Seed-pooled phase MSE against $\arg \xi(\frac{1}{2} + it)$ is 2.1708 on the ξ-zero set and 0.1195 on the off-zero set, a ratio of 18.16× and a Welch $p = 1.18 \times 10^{-15}$. Per-seed ratios are 29.5, 31.9, 9.5 with $p \in \{4.5, 3.7, 5.5\} \times 10^{-6}$.

- **Arithmetic read-out, same weights.** The identical training run, read out by the old arithmetic mean, collapses to zero/non-zero means 2.410 vs. 1.849, ratio 1.30, pooled $p = 0.046$. The circular vs. arithmetic gap is therefore not a training difference but a measurement difference.

Both hypotheses of Remark 4.2 are confirmed: the loss-weight imbalance (ii) and the arithmetic-mean read-out (iii) each suffice to hide the signal, and together they fully explain the negative outcomes of Remark 3.2 and of the first supervised control. With both closed, the supervisory chain $\arg \xi(\frac{1}{2} + it) \mapsto \mathcal{L}_{\text{tgt}} \mapsto \hat{\Psi}$ is *intact*: the off-zero MSE of 0.12 shows that the network learns $\arg \xi$ essentially perfectly on the smooth branch.

The residual on-zero MSE of $2.17 \approx (\pi/2)^2$ has a clean interpretation. On the critical line ξ is real; numerically $|\operatorname{Im} \xi|/|\operatorname{Re} \xi| \sim 10^{-53}$ in our cache, so $\arg \xi$ is effectively a $\{0, \pi\}$-valued step function switching sign at every Riemann zero. The network output $\hat{\Psi}(t)$ is a smooth $\mathbb{R}$-valued field and cannot represent such an instantaneous $\pm\pi$ jump; its best circular-mean fit in a neighbourhood of a zero is, up to a constant, the halfway value, contributing exactly $(\pi/2)^2$ to the MSE. The remaining defect is therefore neither in the target nor in the read-out but in the *output representation* of the discriminant network, and is addressed in Remarks 3.2 and 4.2 only by raising L; a proper fix requires either a smoothed zero-crossing target (e.g.

$\mathrm{sign}(Z(t)) \cdot e^{-|Z(t)|^2/\sigma^2}$) or a head that predicts $\xi(\frac{1}{2}+it)$ as a complex vector and reads out arg post-hoc.

Smooth zero-crossing confirmation (2026-04-10). We tested two smooth surrogates that encode the same sign information as $\arg\xi$ without the $\pm\pi$ discontinuity, under the same isolation setup ($\lambda_{\mathrm{res}} = \lambda_{\mathrm{curv}} = \lambda_{\mathrm{pqo}} = 0$, `in_features`= 128, three seeds):

- **Bump target** $\Psi_{\mathrm{bump}}(t) = \mathrm{sign}(\mathrm{Re}\,\xi) \cdot e^{-(\mathrm{Re}\,\xi)^2/\sigma^2}$, $\sigma = \mathrm{median}(|\mathrm{Re}\,\xi|)$: zero MSE $= 0.935$, off-zero MSE $= 0.693$, ratio $1.35\times$ (gap reduction 92.6%).
- **Normalised target** $\Psi_{\mathrm{norm}}(t) = \mathrm{Re}\,\xi(\frac{1}{2}+it)/\max|\mathrm{Re}\,\xi|$: zero MSE $= 1.8 \times 10^{-5}$, off-zero MSE $= 5.8 \times 10^{-4}$, ratio $0.03\times$ (gap reduction 99.8%).

The normalised target eliminates the zero/non-zero gap entirely — in fact *inverting* it, with zeros easier than non-zeros. This confirms that the $\pm\pi$ discontinuity was the *sole* bottleneck: the network architecture and training pipeline have no residual deficiency once the target is smooth. The analytic content of the programme survives; only the $\{0, \pi\}$ encoding of the target needs replacement. (Script `scripts/smooth_zero_crossing.py`; output `smooth_zero_crossing.{json,txt}`.)

We record the isolation and smooth-target experiments here as evidence that the analytic content of the programme — the identification $\Psi(t) = \arg\xi(\frac{1}{2}+it)$ of Equation (44) — is the correct self-dual target; the smooth-target variant resolves the remaining representability barrier.

4.4 Connection to the Riemann–Siegel ϑ Function

The classical phase analysis of ζ on the critical line uses the Riemann–Siegel theta function:

$$\vartheta(t) = \arg\Gamma\left(\frac{1/4 + it/2}{1}\right) - \frac{t}{2}\log\pi \;\approx\; \frac{t}{2}\log\frac{t}{2\pi} - \frac{t}{2} - \frac{\pi}{8} + O(t^{-1}). \tag{45}$$

The Hardy Z-function is defined as:

$$Z(t) = e^{i\vartheta(t)}\zeta(\tfrac{1}{2}+it), \tag{46}$$

which is real for real t. Its sign changes mark the zeros of ζ on the critical line. The Gram points g_n are defined by $\vartheta(g_n) = n\pi$, providing a phase-quantised grid along the t-axis.

Definition 4.4 (Gram Points)**.** The n-th Gram point g_n satisfies:

$$\vartheta(g_n) = n\pi, \quad n \in \mathbb{Z}. \tag{47}$$

Between consecutive Gram points, the phase $\vartheta(t)$ advances by exactly π. By the argument principle, the number of zeros of $\zeta(\frac{1}{2}+it)$ up to height T is:

$$N(T) = \frac{1}{\pi}\Big[\vartheta(T) + \mathrm{Im}\log\zeta(\tfrac{1}{2}+iT)\Big] + O(1). \tag{48}$$

Remark 4.4 (QROP-Net Parallel: Phase Grid)**.** The Gram points $\{\vartheta(g_n) = n\pi\}$ form a discrete phase lattice on the critical line, directly analogous to the 2^{d_u} discrete phase levels on the SLM in QROP-Net. In both systems, the fundamental question is: *does the continuous signal (the optical wavefront, the ξ-function) align with the prescribed discrete phase grid?*

5 Gate Structure: Phase Resonance as Prime Detection

We now formalise the two-gate criterion for identifying "resonant" points on or near the critical line.

5.1 Gate A: Phase Quantisation (Alignment)

Definition 5.1 (Gate A — Phase Quantisation Condition)**.** Consider a path in the complex s-plane from a reference point $(1/2, t_{\text{ref}})$ to (σ, t). The total accumulated phase change along this path is:

$$\Delta v(\sigma, t) = \int_{t_{\text{ref}}}^{t} v_t(\tfrac{1}{2}, \tau)\, d\tau + \int_{1/2}^{\sigma} v_\sigma(\sigma', t)\, d\sigma'. \tag{49}$$

Gate A requires this phase accumulation to be an integer multiple of π:

$$\boxed{\mathsf{G}_{\text{A}} : \quad \Delta v(\sigma, t) = m\pi, \quad m \in \mathbb{Z}.} \tag{50}$$

On the critical line ($\sigma = 1/2$), the path reduces to the t-direction only, and Gate A becomes:

$$\Delta v(\tfrac{1}{2}, t) = \vartheta(t) - \vartheta(t_{\text{ref}}) + \arg\zeta(\tfrac{1}{2} + it) - \arg\zeta(\tfrac{1}{2} + it_{\text{ref}}) = m\pi, \tag{51}$$

which, for $t_{\text{ref}} = 0$ and $\vartheta(0) = 0$, reduces to the Gram point condition $\vartheta(t) = m\pi$ (up to the $\arg\zeta$ correction).

Remark 5.1 (QROP-Net Parallel: $\sin^2$ Quantisation)**.** Gate A is the *exact analytic number theory counterpart* of the QROP-Net quantisation loss $\mathcal{L}_{\text{quant}}$ (Equation (20)). In QROP-Net, the $\sin^2$ penalty drives continuous phase to the nearest discrete level; Gate A selects points where the continuous phase $v(\sigma, t)$ is quantised to integer multiples of π. The phase quantisation operator (Equation (30)) with $L = 2$ gives $\mathcal{P}(v; 2) = \sin^2(v) = 0$ at $v = m\pi$, which is exactly Gate A.

5.2 Gate B: Phase Stationarity (Curvature)

Definition 5.2 (Gate B — Phase Stationarity Condition)**.** The phase is *stationary* in the σ-direction:

$$\boxed{\mathsf{G}_{\text{B}} : \quad v_\sigma(\sigma, t) = 0.} \tag{52}$$

By the Cauchy–Riemann relations (Equation (40)), this is equivalent to:

$$\operatorname{Im}\left(\frac{\xi'(\sigma + it)}{\xi(\sigma + it)}\right) = 0, \tag{53}$$

and also to the amplitude extremum condition $u_t(\sigma, t) = 0$, *i.e.*,

$$\frac{\partial}{\partial t}|\xi(\sigma + it)| = 0. \tag{54}$$

Gate B selects points where the amplitude $A(\sigma, t)$ has a local maximum (or saddle) in the t-direction.

Remark 5.2 (Physical Interpretation: Resonance Peak)**.** In scattering theory [12], $\xi(s)$ can be interpreted as a scattering amplitude, with its nontrivial zeros corresponding to resonances. At a resonance, the scattering amplitude reaches a *local maximum* (Gate B: $u_t = 0$) and the phase shift passes through a π-increment (Gate A: $\Delta v = m\pi$). This is precisely the condition for constructive interference in a resonant cavity: the wave "lingers" at the resonance, building up amplitude while maintaining phase coherence.

Remark 5.3 (QROP-Net Parallel: Fidelity Loss)**.** Gate B requires the amplitude $A(t)$ to be stationary at the point of interest. In QROP-Net, the fidelity loss $\mathcal{L}_{\text{fid}}$ (Equation (19)) compares the reconstructed amplitude pattern with the target. Minimising $\mathcal{L}_{\text{fid}}$ drives the system toward amplitude stationarity: at convergence, the predicted intensity matches the target, meaning $\partial \hat{I}/\partial(\text{params}) = 0$. Both systems enforce amplitude stability as a necessary condition for correct signal recovery.

5.3 Combined Discriminant: Resonant Points

Definition 5.3 (Resonant Prime Selection)**.** A point (σ^*, t^*) in the s-plane is a *resonant point* if both gates are simultaneously satisfied:

$$\boxed{\pi_{\text{s}}(\sigma^*, t^*) = 1 \quad \Longleftrightarrow \quad \begin{cases} \mathsf{G}_{\text{A}}: & \Delta v(\sigma^*, t^*) = m\pi \\ \mathsf{G}_{\text{B}}: & v_\sigma(\sigma^*, t^*) = 0 \end{cases}} \tag{55}$$

If $\sigma^* = \frac{1}{2}$ (on the critical line), the resonant point corresponds to a zero of $\zeta(s)$ or a near-zero. An integer n is selected as a *resonant prime* if the phase alignment condition $\Phi_\xi(t_n) = \Psi(n)$ holds at a resonant point t_n.

The simultaneous satisfaction of both conditions is the defining characteristic:

- Gate A alone selects too many points (all Gram points, including those between zeros).
- Gate B alone selects amplitude extrema that may not be phase-aligned.
- Their *intersection* yields candidates that are both phase-quantised and amplitude-extremal—the signatures of true resonances.

6 The σ-Lift and Error Tolerance

6.1 Deviation from the Critical Line

Definition 6.1 (σ-Lift)**.** For a resonant point (σ^*, t^*) satisfying both gates, the σ-lift is:

$$\ell := \sigma^* - \frac{1}{2}. \tag{56}$$

The Riemann Hypothesis is equivalent to the statement that $\ell = 0$ for all nontrivial resonant points.

Proposition 6.1 (Error Threshold Analogy)**.** *The σ-lift ℓ plays the same structural role as the LWE error $|e|_{\mathbb{Z}_q}$ in Kyber:*

	QROP-Net	***Critical Line***
Variable	$\lvert e\rvert_{\mathbb{Z}_q}$	$\ell = \sigma^* - 1/2$
Threshold	$\lceil q/4 \rfloor = 1920$	$\ell = 0$
Success	$\lvert e\rvert_{\mathbb{Z}_q} < 1920$	$\ell = 0$ *(RH)*
Failure mode	DFR $= 1$	*Zero off critical line*
Sensitivity	$\Delta d = 0.1\,mm$	$\ell \neq 0$

In both systems, there exists a ***hard boundary*** *beyond which the signal is irrecoverably lost. The "error" in one system is a literal decryption failure; in the other, it is a migration of the zero away from the line of symmetry.*

6.2 Scattering Interpretation and Unitarity

Khare [12] proposed interpreting $\zeta(s)$ as a scattering amplitude, with the nontrivial zeros as resonances. LeClair and Mussardo [13] formalised this by treating the zeros as *quantised energies of a scattering system*, deriving a Bethe-ansatz/phase-quantisation approach to locate them.

In their framework, a zero at $s = \sigma^* + it^*$ with $\sigma^* \neq \frac{1}{2}$ corresponds to a *non-unitary* scattering matrix: the system has absorption or leakage. The σ-lift ℓ quantifies the *degree of non-unitarity*:

$$|\ell| = |\sigma^* - \tfrac{1}{2}| \propto \text{“absorption coefficient”}, \tag{57}$$

where $\ell = 0$ corresponds to a unitary (lossless) S-matrix and $\ell \neq 0$ to a dissipative channel.

Remark 6.1 (QROP-Net Parallel: Channel Loss)**.** In QROP-Net, the optical channel introduces noise (Poisson, read noise, ADC quantisation) that degrades the phase signal. The "channel quality" is directly measured by the LWE error: higher noise $\to$ larger $|e|_{\mathbb{Z}_q}$ $\to$ closer to the DFR threshold. Similarly, $|\ell| > 0$ indicates "noise" in the analytic continuation—the zero is not exactly on the symmetry axis. Both systems measure signal degradation by the distance from an ideal reference (the lattice point, the critical line).

6.3 Berry Phase and Holonomy

The phase winding around a zero of $\xi(s)$ provides a topological invariant.

Proposition 6.2 (Phase Winding Number)**.** *Let $\mathcal{C}$ be a closed contour in the s-plane encircling exactly N zeros of $\xi(s)$. Then:*

$$\oint_{\mathcal{C}} \nabla v \cdot d\mathbf{s} = 2\pi N, \tag{58}$$

by the argument principle. This is a statement of holonomy: the total phase change around the loop is 2π per enclosed zero, directly analogous to the Berry phase [9] in quantum mechanics, where the geometric phase accumulated over a cyclic path is 2π times the Chern number.

Gate A's condition $\Delta v = m\pi$ along an open path from the reference point to (σ, t) is a *partial holonomy* condition: the path accumulates m half-turns, fixing a *topological phase index* m for each resonant point.

Remark 6.2 (QROP-Net Parallel: Phase Wrapping)**.** In QROP-Net, the 2-channel $(\mathrm{Re}, \mathrm{Im})$ output was specifically designed to avoid the phase-wrapping discontinuity at $0/2\pi$. The same topological obstruction appears in the ξ-function: $v(\sigma, t)$ has branch cuts and 2π jumps at zeros. The winding number (Equation (58)) counts these jumps, just as QROP-Net's circular statistics (Equation (25)) handle the $0/2\pi$ boundary correctly.

7 Multi-Group Zero Detection: The Polyadic Frame

The numerical implementation of the gate conditions requires robust evaluation of the logarithmic derivative ξ'/ξ on the critical line, where ξ has zeros (causing ξ'/ξ to diverge). The *polyadic frame* approach resolves this by using multiple directional channels to average out the singularity.

7.1 Target Functions and Evaluation Points

The target function f is unified across the Riemann ξ and Dirichlet L-functions:

$$\xi(s) = \frac{1}{2} s(s-1)\pi^{-s/2}\Gamma(s/2)\zeta(s), \tag{59}$$

$$\Lambda(s, \chi) = \left(\frac{q}{\pi}\right)^{(s+\epsilon)/2} \Gamma\left(\frac{s+\epsilon}{2}\right) L(s, \chi), \tag{60}$$

where q is the modulus, $\epsilon \in \{0, 1\}$ is the parity indicator of the Dirichlet character χ, and all evaluations are on the critical line $s(t) = \frac{1}{2} + it$.

The logarithmic derivative connects directly to primes via the Euler product:

$$\frac{f'(s)}{f(s)} = -\sum_{\rho} \frac{1}{s-\rho} + (\text{smooth part}), \tag{61}$$

where the sum runs over the nontrivial zeros ρ, and the dominant contributions come from

nearby zeros. This is directly related to the Euler product representation:

$$\zeta(s) = \prod_{p \text{ prime}} (1 - p^{-s})^{-1}, \qquad L(s, \chi) = \prod_p (1 - \chi(p) p^{-s})^{-1}, \tag{62}$$

from which:

$$-\frac{f'(s)}{f(s)} \approx \sum_{n \geq 1} \Lambda(n) \chi(n)\, n^{-s}, \tag{63}$$

where $\Lambda(n)$ is the von Mangoldt function ($\Lambda(n) = \log p$ if $n = p^k$, else 0). The oscillatory sums over primes directly produce the vibrations that Gate A quantises.

7.2 Directional Channels

Rather than evaluating f'/f along a single direction (which is sensitive to noise and aliasing), we introduce m directional channels. Each channel k has a basis vector:

$$\mathbf{e}_k = (\cos\theta_k,\ \lambda_k \sin\theta_k), \tag{64}$$

where θ_k is the channel angle and λ_k is an anisotropic scaling factor. The recommended 3-channel standard set is:

$$(\theta, \lambda, w) = (-30^\circ, 1.10, 0.40),\ (15^\circ, 0.85, 0.35),\ (60^\circ, 1.00, 0.25), \tag{65}$$

with channel weights w_k normalised to $\sum w_k = 1$.

An anisotropic scaling matrix adapts to the local geometry:

$$M(t) = \operatorname{diag}\left(\frac{1}{\mu(t)}, \frac{1}{\nu(t)}\right), \tag{66}$$

where $\mu(t), \nu(t) > 0$ are smooth profile functions. The actual directional displacement for step $h > 0$ is:

$$\delta_k(t) = h\,(v_{k,x} + i v_{k,y}), \quad \mathbf{v}_k = R(\theta(t))\, M(t)\, \mathbf{e}_k, \tag{67}$$

where $R(\theta)$ is a rotation matrix and $\theta(t)$ varies smoothly around t_0.

7.3 Central Difference Log-Derivative

The effective log-derivative along channel k is computed via central differences:

$$D_k \xi(s) \approx \frac{\xi(s + \delta_k) - \xi(s - \delta_k)}{2h\, \xi(s)}. \tag{68}$$

The *weighted multi-channel log-derivative* is:

$$\hat{L}(t) = \sum_k w_k \cdot \frac{f(s + \delta_k) - f(s - \delta_k)}{2h\, f(s)}, \tag{69}$$

where the channel averaging reduces directional bias and improves stability near zeros. Since f is analytic, the central difference has error $O(h^2)$. We decompose $\hat{L}(t) = v_t(t) + i\, v_\sigma(t)$ into

its real and imaginary parts.

7.4 Gauge ODE for Phase Correction

The real part $v_t(t)$ contains a slow drift that must be removed before the final discriminant can be applied. We define a gauge ODE:

$$\tau\dot{\psi} + \psi = v_t, \quad \dot{\phi} = \psi, \tag{70}$$

with initial conditions $\phi(t_0) = 0$, $\psi(t_0) = v_t(t_0)$. Here $\tau > 0$ is the gauge damping constant (default: $\tau = 0.04$). The ODE integrates v_t while smoothing out high-frequency noise. The numerical scheme uses a stable trapezoidal discretisation:

$$\alpha_j = \frac{\Delta t_j}{\tau + \Delta t_j}, \tag{71}$$

$$\psi_j = \psi_{j-1} + \alpha_j(v_{t,j} - \psi_{j-1}), \tag{72}$$

$$\phi_j = \phi_{j-1} + \tfrac{1}{2}(\psi_j + \psi_{j-1})\,\Delta t_j. \tag{73}$$

7.5 Final Discriminant Function

The gauge-corrected final discriminant combines both gates:

$$\boxed{F_2(t) = \mathrm{Im}\Big\{e^{-i\phi(t)}\Big(\hat{L}(t) - \dot{\phi}(t)\Big)\Big\} = 0.} \tag{74}$$

At a zero t^* of F_2:

- The imaginary part vanishes, meaning the phase-corrected log-derivative is aligned with the accumulated gauge phase → **Gate A** (phase quantisation) is satisfied.
- The sign change of F_2 through zero indicates a phase crossing → **Gate B** (stationarity) is locally satisfied.

Remark 7.1 (QROP-Net Parallel: Blind Self-Calibration)**.** The gauge ODE (Equation (70)) self-corrects the phase drift *without knowing the true zero locations.* This is structurally identical to QROP-Net's blind TTO (Equation (29)), which corrects lens misalignment *without knowing the true ciphertext.* Both systems:

1. Use physics/mathematics-based constraints (not ground truth),
2. Converge in a small number of steps (15 for TTO, $O(N)$ grid points for F_2),
3. Achieve self-calibration through internal consistency of the signal structure.

7.6 Numerical Verification and Convergence

The discriminant $F_2(t) = 0$ is searched numerically on a uniform grid $[T - W, T + W]$ with N points, where $W = 0.002$–0.005. The algorithm:

Algorithm 1: Polyadic Zero Detection

Input: Target f, band centre T, window W, channels $\{(\theta_k, \lambda_k, w_k)\}$, step h, gauge τ
Output: Zero candidates $\{t^*\}$

$s_j \leftarrow \frac{1}{2} + it_j$ for uniform grid $t_j \in [T - W, T + W]$
foreach *channel k and grid point j* **do**
  $\delta_k(t_j) \leftarrow h\,(v_{k,x} + iv_{k,y})$ via Equation (67)
  $D_k \leftarrow [f(s_j + \delta_k) - f(s_j - \delta_k)]/(2h\,f(s_j))$
$\hat{L}(t_j) \leftarrow \sum_k w_k\, D_k$
Extract $v_t \leftarrow \mathrm{Re}(\hat{L})$, $v_\sigma \leftarrow \mathrm{Im}(\hat{L})$
Integrate gauge ODE (Equation (71)–73) to get $\phi(t_j)$
$F_2(t_j) \leftarrow \mathrm{Im}\{e^{-i\phi_j}(\hat{L}_j - \dot{\phi}_j)\}$
Detect sign changes of F_2; refine by linear interpolation
Verify invariance: Frame M_0 vs M_2, channel swap, window shift

Property 7.1 (Convergence Rate)**.** Let $h_{\min}$ be the smallest step tested. The convergence indicators are:

$$p_{\mathrm{LD}} = \frac{\log \sup_t |\mathrm{LD}(h) - \mathrm{LD}(h_{\min})|}{\log h} \approx 2.0, \tag{75}$$

$$p_{F_2} = \frac{\log \sup_t |F_2(h) - F_2(h_{\min})|}{\log h} \approx 2.0, \tag{76}$$

confirming $O(h^2)$ convergence as expected for central differences. The zero location t^* is invariant to:

- Frame profile (M_0 flat vs M_2 strong): $|t^*_{M_0} - t^*_{M_2}|/\Delta t_{\mathrm{grid}} < 10^{-3}$.
- Channel configuration (3-channel vs 2-channel): $r_{\mathrm{channel}} < 0.1$.
- Window shift ($\pm\delta$ perturbation): $r_{\mathrm{window}} < 0.2$.

Verified targets include:

Function	**Band**	t^*
$\xi(s)$	$t \approx 141.12$	141.123703
$L(s, \chi_3^{\mathrm{odd}})$	$t \approx 140.61$	140.611723
$L(s, \chi_8^{\mathrm{even}})$	$t \approx 141.08$	141.083569
$L(s, \chi_5^{\mathrm{cmplx}})$	$t \approx 140.97$	140.966280

All results are invariant under the step sweep $h \in \{2 \times 10^{-7}, 10^{-6}, 5 \times 10^{-6}, 10^{-5}\}$.

7.7 Cross-Validation: Three Independent Axes

The zero candidates from $F_2 = 0$ are validated against three independent data sources with fundamentally different error profiles:

7.7.1 Axis A: Partial Euler Product Matching

The partial sum $S_N(t;\chi) = \sum_{n\le N} \Lambda(n)\chi(n)\, n^{-1/2-it}$ approximates $-f'(s)/f(s)$ near $s = \frac{1}{2}+it$ for sufficiently large N. A match between $\hat{L}(t^*)$ and $S_N(t^*;\chi)$ within tolerance confirms that the zero candidate is consistent with the Euler product.

7.7.2 Axis B: Explicit Chebyshev $\psi(x)$ Recovery

The prime-counting function $\psi(x) = \sum_{n\le x} \Lambda(n)$ can be reconstructed from the zeros via the explicit formula:

$$\psi(x) = x - \sum_{\rho} \frac{x^\rho}{\rho} - \log(2\pi) - \frac{1}{2}\log(1-x^{-2}), \tag{77}$$

where ρ runs over the nontrivial zeros of ζ. Inserting the zero heights t^* from $F_2 = 0$ and comparing the reconstructed $\psi(x)$ with the exact step function provides an independent verification.

7.7.3 Axis C: Arithmetic Progression Bias Signal

For Dirichlet L-functions, the distribution of primes in arithmetic progressions $p \equiv a \pmod q$ provides a third signal. The bias term $\sum_\rho x^\rho/\rho$ in the explicit formula for $\psi(x;q,a)$ oscillates with frequencies determined by the zero heights, providing independent confirmation.

8 ■ The F_2-Discriminant as a Hardy-Z Sign-Change Detector

Tier 1. This section identifies the central validated empirical result of the programme: the multi-channel discriminant F_2 introduced through the gauge ODE of Section 7 is, up to an analytic signal construction, exactly the Hardy Z-function sign-change detector. The identification is established in two steps: an analytic statement about the argument of the analytic signal of a real function at its zeros (Proposition 8.1), and a neural corollary explaining why a smooth network approximating $Z(t)$ saturates at $\arg(Z_{\text{out}}) = \pi/2$ at simple zeros (Corollary 8.2).

8.1 Hardy Z and the Analytic-Signal Identity

We recall from Equation (46) that the Hardy Z-function

$$Z(t) = e^{i\vartheta(t)}\,\zeta(\tfrac{1}{2}+it) \tag{78}$$

is real for real t and that its sign changes are in bijection with the non-trivial zeros of ζ on the critical line. This makes Z the natural real-valued signal on which to test a discriminant designed to locate Riemann zeros.

Proposition 8.1 (Phase at zeros of real functions)**.** *Let $f : \mathbb{R} \to \mathbb{R}$ be a real-valued function with a simple zero at t_0, and let $z(t) = f(t) + i\,H[f](t)$ be its analytic signal, where H denotes the Hilbert transform. Then*

$$\arg\big(z(t_0)\big) = \pm\tfrac{\pi}{2},$$

with the sign determined by whether f crosses from negative to positive or vice versa.

Proof. At t_0, $f(t_0) = 0$ so $\operatorname{Re} z(t_0) = 0$, while $H[f](t_0) \neq 0$ generically by Titchmarsh's theorem. Hence $z(t_0)$ is purely imaginary, with sign determined by the direction of the crossing. □

Corollary 8.2 (Neural sign-change regularisation)**.** *A neural network approximating $Z(t)$ as a continuous complex-valued function $Z_{\text{out}}(t)$ necessarily satisfies* $\arg(Z_{\text{out}}(t_0)) \to \pm\pi/2$ *at simple zeros t_0, since the ideal discontinuity* $\arg \xi : 0 \to \pi$ *of Equation (41) is regularised to its midpoint.*

Combining Proposition 8.1 and Corollary 8.2 with the definition of the discriminant $F_2(t) = \operatorname{Im}\{e^{-i\varphi}(\hat{L} - \dot{\varphi})\}$ from Section 7 gives the following identification.

Theorem 8.3 (F_2 = Hardy Z sign-change detector)**.** *Let $Z_{\text{out}}(t)$ be a smooth neural approximation of $Z(t)$ satisfying the gauge ODE $\dot{\varphi} = \hat{L}(t)$ of Equation (70). Then on any interval $I \subset \mathbb{R}$ where the approximation is faithful,*

$$\{\, t \in I : F_2(t) = 0 \,\} \;=\; \{\, t \in I : Z(t) = 0 \,\} \;=\; \{\, t \in I : \zeta(\tfrac{1}{2} + it) = 0 \,\}.$$

Proof. $F_2(t) = \operatorname{Im}\{e^{-i\varphi(t)}(\hat{L}(t) - \dot{\varphi}(t))\} = 0$ iff $\hat{L}(t) - \dot{\varphi}(t)$ is colinear with $e^{i\varphi(t)}$, which by the gauge ODE is equivalent to the analytic-signal argument reaching the midpoint $\pm\pi/2$, i.e. by Proposition 8.1 to a zero of Z. □

8.2 Empirical Validation: 10/10 Zero Detection on $t \in [10, 50]$

Observation 8.4 (F_2 detection, low range)**.** *With Adam ($\eta = 10^{-3}$, 200 epochs), the smooth neural approximation of $Z(t)$ on $t \in [10, 50]$ converges to $|F_2(t_k)| < 0.021$ at all ten LMFDB zeros, with residual loss $\mathcal{L}_{\text{res}} < 10^{-4}$ after epoch 15. The output phases* $\arg(Z_{\text{out}}) \in [1.47, 1.62]$ *fall inside a ± 0.06 window around $\pi/2 \approx 1.5708$, in precise agreement with Corollary 8.2.*

Table 1: Adam + $\cos^2$ PQO detection results on $t \in [10, 50]$, 10 LMFDB zeros, 200 epochs.

t_k	$\lvert F_2 \rvert$	$\sin^2$ PQO	$\cos^2$ PQO	$\arg(Z_{\text{out}})$
14.13	0.021	1.000	**0.000**	1.581
21.02	0.001	1.000	**0.000**	1.563
25.01	0.003	1.000	**0.000**	1.557
30.42	0.019	1.000	**0.000**	1.562
32.94	0.000	1.000	**0.000**	1.562
37.59	0.009	1.000	**0.000**	1.571
40.92	0.000	1.000	**0.000**	1.565
43.33	0.004	1.000	**0.000**	1.573
48.01	0.003	1.000	**0.000**	1.584
49.77	0.005	1.000	**0.000**	1.582

Two observations deserve emphasis. First, the simultaneous satisfaction $|F_2| < 0.021$ and $\cos^2 \text{PQO} < 0.0002$ is a non-trivial joint condition: it says the network has located a point where (i) the analytic signal's imaginary component vanishes and (ii) the quantised phase

attractors of the cyclotomic lattice are satisfied. The fact that both conditions can be met simultaneously on the same ten points is the $\cos^2$ Maslov correction at work. Second, the residual gap $\arg(Z_{\text{out}}) \neq \pi/2$ (at most ≈ 0.01) is an artefact of finite training, not a structural obstruction; the gap narrows monotonically with training time.

8.3 Zero Difficulty Hierarchy and the Lehmer Phenomenon

Not all zeros are equally easy to detect. The ten zeros of Table 1 span nearly two orders of magnitude in $|F_2|$: the easiest zeros converge to $|F_2| \lesssim 10^{-3}$, while the hardest ($t_1 = 14.13$ and $t_6 = 37.59$) saturate around 2×10^{-2}.

Remark 8.1 (Lehmer–Odlyzko connection)**.** Zero pairs where $Z(t)$ nearly grazes the real axis without crossing it — the classical Lehmer phenomenon — correspond to close zero pairs in the GUE spacing statistics of Montgomery–Odlyzko. In our setting, these are precisely the zeros where the analytic signal's phase spends longer time near $\pm\pi/2$. Numerical evidence (78 zeros, $t \in [0, 200]$) indicates that the bundle curvature profile κ at close pairs is qualitatively broader (FWHM ratio ≈ 3.4, p=0.01) with multi-peaked structure.

A high-resolution follow-up confirms the curvature–spacing link via a δ-robustness test: the Spearman rank correlation between the normalised gap g_{norm} and peak curvature $\kappa(t_0 + \delta)$ satisfies $\rho = 0.835$ ($p < 10^{-21}$, $n = 78$) consistently across three offset values $\delta \in \{0.01, 0.03, 0.05\}$ (all $|\rho| > 0.83$). The FWHM–spacing correlation, however, remains weak ($\rho = -0.276$, $p = 0.014$), indicating that spacing modulates peak *height* but not peak *width*. [confirmed 2026-04-14, δ-robust 3/3]

9 ■ Extension to High Zeros: $t \in [100, 200]$

Tier 1. The identification $F_2 =$ Hardy-Z sign-change detector of Theorem 8.3 is a statement about the analytic signal of $Z(t)$, with no intrinsic height dependence. A natural sanity check is whether the numerical pipeline of Section 8.2 continues to produce complete detection when the target range is raised by an order of magnitude. We find that it does.

Observation 9.1 (High-zero detection)**.** *The experiment of Section 8.2 was repeated on $t \in [100, 200]$ with 50 LMFDB zeros (computed via* `mpmath.zetazero`(k) *for $k = 26, \ldots, 79$), 500 grid points, and an enlarged network (hidden width 64, 33,540 parameters). With Adam + $\cos^2$ PQO, $\eta = 10^{-3}$, 200 epochs:*

Table 2: Low-range vs. high-range detection.

Metric	$t \in [10, 50]$, 10 zeros	$t \in [100, 200]$, 50 zeros
Detection rate	10/10	**50/50**
$\max \lvert F_2 \rvert$	0.021	0.060
$\max \cos^2$ PQO	0.0002	0.0011
$\arg(Z_{\text{out}})$ range	$[1.55, 1.58]$	$[1.54, 1.60]$
Training wall time	2 min	19 min

All fifty high zeros satisfy $|F_2| < 0.1$, with forty-eight of them below 0.04. The two hardest — $t = 118.79$ ($|F_2| = 0.060$) and $t = 114.32$ ($|F_2| = 0.058$) — are still comfortably inside the detection threshold. The critical synchronisation transition of Remark 13.1 at ~ 10 epochs and the $\arg(Z_{\text{out}}) \approx \pi/2$ pattern are identical to the low range, confirming that the Maslov correction and the phase-locking mechanism are not artefacts of a particular height window. This is the direct numerical corroboration of Theorem 8.3 at a non-trivial scale.

Remark 9.1 (What the extension does *not* establish)**.** The experiment verifies that the F_2-detector continues to work when its ground truth — the zero list — is supplied. It does *not* verify that the detector would *find* unknown zeros: the training loop uses the LMFDB zero list to locate the grid points on which F_2 is evaluated. The question of self-calibrating, blind zero discovery remains open and is one of the motivations for Part III and the open-problem discussion of Section 15.3.

10 ■ The $|F_2|$ residual: a tentative mesoscopic observable

Tier 2 (with 2026-04-08 partial reinterpretation). While running the high-zero extension of Section 9 we noticed that the post-convergence magnitudes $|F_2(t_k)|$ of the discriminant at individual zeros are not a featureless random scatter but display systematic height and spatial patterns. This section records the original observations (April 2026, 14,441 zeros) and then the 2026-04-08 partial reinterpretation that weakened the per-zero claim while leaving the aggregate scale claim intact. We keep both so that the reader can see how the claim was narrowed.

10.1 Original observations (April 2026)

The $|F_2|$ residual was computed on 14,441 zeros across seven wide intervals $[10, 50]$ through $[5000, 10000]$ and 23 fine intervals covering $[10000, 13000]$. The cumulative detection rate (zeros for which $|F_2|$ is below a uniform threshold) was 99.99%.

Table 3: $|F_2|$ mean residual by height interval, from the original 14,441-zero sweep.

Range	**Zeros**	**Detection rate**	$\overline{\lvert F_2 \rvert}$
$[10, 50]$	10	10/10	0.0209
$[100, 200]$	50	49/50	0.0208
$[200, 500]$	190	190/190	0.0119
$[500, 1000]$	380	380/380	0.0083
$[1000, 2000]$	868	868/868	0.0058
$[2000, 5000]$	3,003	3,003/3,003	0.0020
$[5000, 10000]$	5,622	5,622/5,622	0.0014

Three per-zero claims were originally made from this dataset:

(O1) **Independence from classical statistics.** The per-zero residual $|F_2(t_k)|$ showed Pearson $|r| < 0.09$ and Spearman $|\rho| < 0.09$ against all of $|Z'(t_k)|$, $|Z''(t_k)|$, normalised spacing s_k, the argument function $S(t_k)$, the Riemann–Siegel derivative $\vartheta'(t_k)$, local density, spacing variance, spacing autocorrelation, and the local Z-area $\int_{t_k-0.3}^{t_k+0.3} |Z|$.

(O2) **Spatial clustering.** $|F_2(t_k)|$ was Spearman-correlated with the mean of $|F_2|$ at its six nearest zero neighbours, $\rho = +0.895$ ($p \approx 10^{-141}$), interpreted as evidence for "mesoscopic domains" of different character along the critical line.

(O3) **Long-range autocorrelation.** The lag-Δ autocorrelation $C(\Delta)$ dropped rapidly to ≈ 0.4 within $\Delta \sim 3$ and then remained nearly flat at ≈ 0.33 out to $\Delta = 200$, suggesting a two-component (short-range exponential + long-range plateau) structure.

Taken together these were read as evidence that the gauge ODE acts as "a new lens on the critical line" revealing structure invisible to random matrix theory and to the classical $S(t)$-based tools.

10.2 2026-04-08 partial reinterpretation

Remark 10.1 (Partial reinterpretation of the $|F_2|$ per-zero claim)**.** A model-independence audit performed on 2026-04-08 substantially weakens (O1)–(O3) at the *per-zero* level while leaving the aggregate scale structure of Table 3 intact.

Specifically, retraining the same architecture with different seeds, different random initialisations of the gauge buffer (φ_0, ψ_0), and different shuffles of the training zero list produces $|F_2(t_k)|$ values whose rank order at individual zeros is *not* reproducible across runs. The per-zero Spearman correlation between two independent runs on the same interval drops below $+0.2$ for most intervals, inconsistent with the interpretation of $|F_2(t_k)|$ as an intrinsic property of the zero t_k. In particular, the spatial clustering statistic $\rho_{\text{cluster}} = +0.895$ is partially inflated by a training-time smoothness prior that couples neighbouring zero residuals; the long-range plateau of $C(\Delta)$ at ≈ 0.33 appears in every run but with a training-dependent level.

What *does* survive reinterpretation is the height-dependent **mean scale** of $\overline{|F_2|}$ across the intervals of Table 3: the monotone decay from $\overline{|F_2|} \approx 0.021$ at $t \sim 50$ to $\overline{|F_2|} \approx 0.0014$ at $t \sim 10{,}000$ is reproduced across all seeds up to $\sim$10% relative variation, and its shape is consistent with a power law $\overline{|F_2|} \sim t^{-\alpha}$ with $\alpha \in [0.4, 0.5]$.

We therefore *retain* Table 3 and the mean-scale claim as Tier 2 evidence of a genuine height dependence in the detector's residual, and we *downgrade* the per-zero claims (O1)–(O3) to "suggestive but model-contaminated": the patterns are real features of the trained model, but the inference from model to underlying arithmetic of the critical line is blocked by the observed non-reproducibility at the per-zero level. The section is retained in this paper precisely because the narrowing is itself an informative diagnostic; see also Section 15.3 on the role of such retraction records.

10.3 2026-04-10 follow-up: the "hard cluster" was a basis-bandwidth artifact

Between 2026-04-09 and 2026-04-10 a ten-experiment cascade (Lehmer regression, five-seed ensemble, hidden-width scaling, sample-density scaling, residual spectrum, cluster physics, input-basis scaling, and a convergence extension) was run against a residual difficulty cluster inside the $|F_2|$ dataset. Specifically, in the trained model (`in_features`= 64, `hidden`= 64,

seed= 42) over the band $t \in [100, 200]$, five zeros in $t \in [165, 173.5]$ carried mean $|F_2| \approx 0.036$, roughly 1.8× the band mean 0.020. The cascade eliminated four candidate explanations and identified a fifth.

Eliminated hypotheses.

(E1) *Lehmer-type near pairs.* A multivariate regression of $|F_2(t_k)|$ on Riemann–von Mangoldt normalised gaps $g_{\min}^{\text{norm}}(t_k)$ across all 15,422 zeros in the 51 training bands, and a separate `mpmath`-based high-t scan, both gave $|r| < 0.03$. The zero with the tightest normalised gap in the cluster band, $t = 121.37$ with $g_{\min}^{\text{norm}} = 0.027$, is in fact the *easiest* zero in its band.

(E2) *Multi-basin seed dependence.* Five-seed ensemble training with identical hyperparameters gave $\text{corr}(\text{mean}, \text{std}) = +0.70$ across zeros, but zero ensemble gain at the hard zeros ($|F_2|^{\text{ensemble}}/|F_2|^{\text{seed-mean}} \approx 1.00$), ruling out a many-local-minima reading of the cluster.

(E3) *Hidden-layer capacity.* Scaling `hidden` $\in \{32, 64, 128\}$ at fixed `in_features`= 64 produced a *monotone opposite* trend: global $\overline{|F_2|}$ went $0.019 \to 0.022 \to 0.030$, so larger hidden width did not help the cluster.

(E4) *Training-sample density.* Increasing `num_points` from 500 to 1500 improved global $\overline{|F_2|}$ by 32% and cluster-max $|F_2|$ by 57%; at $t = 173.41$ the improvement was 65%. This was initially read as a "sampling resolution" cure but turned out to be a proxy for the basis-bandwidth bottleneck identified below.

Identified cause: input-basis bandwidth. The input embedding lifts the scalar t to a Fourier feature vector

$$t \longmapsto \left(\tfrac{t-\mu}{\sigma},\ \sin(\omega_1 t), \cos(\omega_1 t), \ldots, \sin(\omega_K t), \cos(\omega_K t)\right),$$

with $\omega_k = 2\pi k/(t_{\max} - t_{\min})$ and $K = (\texttt{in_features} - 1)/2$. For the band $[100, 200]$ at `in_features`= 64 this gives $K = 31$ and a hard representational ceiling

$$\omega_{\max}^{(64)} = \frac{2\pi \cdot 31}{100} \approx 1.948 \text{ rad/unit } t, \qquad \text{minimum period } \approx 3.23 \text{ in } t.$$

Inside the cluster, the zeros at $t \approx 169.095$ and $t \approx 169.912$ form a pair of separation $\Delta t \approx 0.82$. Representing two sign changes across this separation requires basis content at

$$\omega \gtrsim \frac{\pi}{\Delta t} \approx 3.83 \text{ rad/unit } t, \tag{79}$$

which is $\approx 1.97\times$ the `in_features`= 64 ceiling. We now state this observation as a falsifiable proposition.

Proposition 10.1 (Input-basis bandwidth necessary condition)**.** *Let $\mathcal{F}_K^{\text{lin}}$ denote the class of functions on a band $[t_{\min}, t_{\max}]$ of width $W = t_{\max} - t_{\min}$ that are* linear *in the Fourier feature map $t \mapsto (\sin \omega_k t, \cos \omega_k t)_{k=1}^K$ with $\omega_k = 2\pi k/W$. For any pair of ordinates $t_a, t_b \in [t_{\min}, t_{\max}]$*

with $0 < \Delta t := |t_a - t_b|$*, the following sampling-theoretic bound holds: if*

$$K \;<\; \frac{W}{2\,\Delta t}, \tag{80}$$

then no $f \in \mathcal{F}_K^{\mathrm{lin}}$ *with* $f(t_a) = f(t_b) = 0$ *can have* $f'(t_a) \cdot f'(t_b) < 0$ *while being orthogonal (in* $L^2[t_{\min}, t_{\max}]$*) to the null function. In particular, two sign-changing zeros of* $f \in \mathcal{F}_K^{\mathrm{lin}}$ *at separation* Δt *force* $K \geq W/(2\Delta t)$*.*

Sketch. A function in $\mathcal{F}_K^{\mathrm{lin}}$ is a real trigonometric polynomial of degree $\leq K$ on $[0, W]$ (after recentring). By the Nyquist–Shannon sampling theorem applied to its $2K$-dimensional Fourier basis, its minimum period between consecutive sign changes is bounded below by $W/(2K)$. Two sign changes at separation $\Delta t < W/(2K)$ therefore cannot be realised in $\mathcal{F}_K^{\mathrm{lin}}$, which is (80). □

Remark 10.2 (Nonlinear extension and the intermodulation gap)**.** The network in question is *not* in $\mathcal{F}_K^{\mathrm{lin}}$: the three ResonantBlock layers compose multiplications, ridge nonlinearities, and gauge ODE integration with the Fourier features. In principle, products of the form $\sin(\omega_j t) \cdot \sin(\omega_k t) = \frac{1}{2}[\cos((\omega_j - \omega_k)t) - \cos((\omega_j + \omega_k)t)]$ and higher-order intermodulation cascades produce components up to $\omega_{\max}^{\mathrm{eff}} = L \cdot \omega_{\max}$ where L is the network depth, formally lifting the bandwidth ceiling. Our empirical data (Table 4) show a sharp failure at exactly the linear-basis ceiling (80), *not* at the nonlinearly-lifted ceiling $L \cdot \omega_{\max}$.

This gap is consistent with the *spectral bias* literature for neural networks: under finite training budgets, MLPs preferentially learn the lowest-frequency components first and only activate intermodulation cascades at training times that diverge with frequency [21, 22], a phenomenon derivable from Neural Tangent Kernel eigenvalue decay [20]. Within such a finite-budget regime, the *effective* bandwidth of the trained network remains bounded above by a multiple of the linear-basis ceiling, not by the formal multiplicative lift $L \cdot \omega_{\max}$. A rigorous statement of the form "$\omega_{\max}^{\mathrm{eff}}(E) \leq c(E) \cdot \omega_{\max}^{\mathrm{basis}}$ with $c(E) = O(1)$ as $E \to E_{\max}$ (finite training budget)" for the ResonantBlock architecture is not derived here and is left as an open problem; what the experiment shows is that *empirically* the ceiling behaves as in Proposition 10.1 rather than as $L \cdot \omega_{\max}$.

For the band $[100, 200]$ and the $t \approx 169$ pair, (80) gives $K \geq 61$, equivalently `in_features` $\geq$ 123.

Riemann–von Mangoldt substitution. The condition (80) acquires analytic-number-theoretic content when combined with the Riemann–von Mangoldt local zero density $N'(t) \sim (2\pi)^{-1} \log(t/2\pi)$, which gives mean zero gap

$$\langle \Delta t \rangle(t) \;\sim\; \frac{2\pi}{\log(t/2\pi)}. \tag{81}$$

Assuming a Lehmer-compression factor $c_{\mathrm{L}} \in (0, 1]$ so that the smallest gap in a typical band of width W at height t satisfies $\Delta t_{\min} \gtrsim c_{\mathrm{L}} \cdot \langle \Delta t \rangle(t)$, substitution into (80) yields the

height-dependent bandwidth requirement

$$K(t, W) \;\geq\; \frac{W\,\log(t/2\pi)}{4\pi\, c_{\mathrm{L}}}, \tag{82}$$

or equivalently `in_features` $\geq (W\log(t/2\pi))/(2\pi\, c_{\mathrm{L}}) + 1$. For $W = 100$, $t \approx 150$, and $c_{\mathrm{L}} \approx 1/3$ (a Lehmer-type cluster), this gives $K \gtrsim 82$, `in_features` $\gtrsim 165$, bracketing the empirical transition at `in_features` $\in [64, 128]$. Equation (82) is an architectural–analytic bridge: the Fourier-feature capacity required to detect ξ-zeros in a band scales *logarithmically* with height, matching the $N'(t)$ law. In particular, the apparent improvement of $\overline{|F_2|}$ with height in Table 3 is *compatible* with a fixed `in_features` satisfying (82) only because the high-t bands in the original sweep were run with `num_points` $= \max(200, 10\, n_{\mathrm{zeros}})$ auto-scaling, which preserved the dimensionless ratio $K/(W\log t)$ and therefore kept (82) implicitly satisfied; this connects the height-scaling table to the architectural condition rather than to any analytic property of ξ itself.

Confirmation via basis-scaling. Scaling `in_features`$\in \{32, 64, 128, 256\}$ at fixed `hidden`$=$ 64, seed$=$ 42, `num_points`$=$ 500, training budget 100 epochs gave:

Table 4: Cluster vs. non-cluster $\overline{|F_2|}$ at fixed training budget. The 5-seed ensemble row (seeds $42, 7, 123, 2026, 314$) replaces the single-run claim; cluster improvement is 25% on average but with high seed variance.

`in_feat`	$\omega_{\max}$	global	cluster	other	ratio
32	0.942	0.0140	0.0140	0.0140	1.00×
64	1.948	0.0218	0.0361	0.0202	1.79×
128 (seed 42)	3.958	0.0164	0.0150	0.0166	0.90×
128 (5-seed)	**3.958**	**0.0207**	**0.0269±0.012**	**0.0200**	**1.42 ± 0.73**×
256	7.980	0.0214	0.0581	0.0173	3.36×

Two features stand out. First, the cluster excess at `in_features`$=$ 64 *disappears* at `in_features`$=$ 128 in the best seed, but a five-seed ensemble (seeds $42, 7, 123, 2026, 314$) reveals substantial variability: the cluster-to-non-cluster ratio averages 1.42 ± 0.73 (range 0.61–2.44×), with 3/5 seeds achieving ratio < 1.2 and 2/5 showing a *persisting* or worsened cluster. Per-zero coefficient of variation in the cluster ranges from 46% to 106% across seeds. Thus the bandwidth condition (80) is a *necessary* condition for cluster resolution but not a *sufficient* one: training dynamics (seed-dependent initialisation) retain significant influence even after the Nyquist ceiling is cleared. Averaging over the ensemble, cluster $\overline{|F_2|}$ improves by 25% relative to the `in_features`$=$ 64 baseline (0.0269 vs. 0.0361), confirming the bandwidth mechanism at the population level while excluding any single-run claim of complete resolution. Second, `in_features`$=$ 256 fails to converge in 100 epochs and presents a *new*, unrelated difficulty cluster.

Isolation of the over-parameterisation artefact. Extending training to 300 epochs with 50-epoch best-validation checkpoints separates the bandwidth effect from the optimisation effect. For `in_features`$=$ 128, the best validation loss settles at 0.0550 by epoch 143 and stays flat; the cluster ratio improves to 0.75× and $\overline{|F_2|}^{\mathrm{cluster}}$ drops to 0.0106 versus $\overline{|F_2|}^{\mathrm{other}} = 0.0142$. For `in_features`$=$ 256, the best validation loss stays at 0.224 throughout

300 epochs, roughly 4× worse than the `in_features`= 128 minimum; global $\overline{|F_2|}$ improves from 0.0214 to 0.0099, but the cluster ratio stays at 3.87×. The `in_features`= 256 failure is therefore an independent over-parameterisation pathology (the Adam-with-gradient-clip training budget is insufficient for the enlarged parameter space), not a bandwidth problem; it does not undermine the bandwidth interpretation at `in_features` $\in \{32, 64, 128\}$.

Consequence for the 2026-04-08 reinterpretation. This experiment closes the second open question of Section 10.4: it exhibits a concrete, *model-dependent* per-zero difficulty (the cluster in $[165, 173.5]$) that is not an intrinsic property of the zeros but a property of the embedding $\{\omega_k\}_{k=1}^{K}$. Any attempt to read per-zero $|F_2|$ values as observables on the critical line must therefore first verify that the input-basis bandwidth condition (80) is satisfied for the smallest zero pair in the band under study. The mean-scale claim of Table 3 is unaffected, since it averages over bands in which the condition is (implicitly, via the auto-scaling rule `num_points`= $\max(200, 10 \cdot n_{\text{zeros}})$) met or approximately met.

10.4 What remains open

After the 2026-04-10 follow-up the following questions remain open:

1. **Height scaling law [partially resolved 2026-04-12].** Blind extrapolation recall remains stable at 97.5−98.8% from t∼200 to t∼1100 (3 seeds, $K \in \{128, 256\}$), with precision *improving* at higher t. This rules out gross degradation. The fine-grained question—whether $\overline{|F_2|}$ follows a genuine power law whose exponent matches $\log\log t$ corrections to $N(t)$—still requires a continuous sweep with error bars over disjoint seed families.

2. **Bandwidth-clean per-zero observable.** A five-seed ensemble at `in_features`= 128 confirms that per-zero $|F_2|$ values have coefficient of variation 46–106% across seeds even after (80) is satisfied. Individual per-zero values are therefore *not* reproducible in any naive sense. The *ensemble mean* cluster $\overline{|F_2|}$ improves by 25% relative to the bandwidth-deficient baseline, confirming the population-level mechanism. Whether a tighter invariant can be extracted (e.g. the *minimum* of $|F_2(t_k)|$ over a bootstrap ensemble, interpreted as the "achievable residual") remains open.

3. **Over-parameterisation failure mode.** The `in_features`= 256 failure to converge within 300 epochs (best validation 0.224 versus 0.055 at `in_features`= 128) exhibits a "new" cluster at the same $t \in [165, 173.5]$ that cannot be bandwidth in origin. Whether this is (a) an Adam-plus-gradient-clip training-schedule artefact, (b) a genuine over-parameterisation pathology that revives the same cluster through a different mechanism, or (c) a seed accident remains unresolved.

Part III

Failed Directions and Their Diagnoses

This part collects, in compact form, the empirical programmes that did not succeed. Each section is labelled ■ (Tier 3); each opens with a one-paragraph description of the programme, followed by a diagnostic chain, and closes with the lesson extracted. The content is deliberately dense — typically one to two pages per failed direction — so as not to crowd out the validated and conjectural material while retaining the full diagnosis. A reader in a hurry may skim Part III; a reader planning a follow-up experiment should read it carefully.

11 ■ Cyclotomic $\cos^2$ PQO training ($L = 2, 4$): Ψ-decoupling and the isolation cascade

Programme. Train a three-layer resonant discriminator network to separate Riemann zeros from non-zero ordinates on $t \in [100, 200]$ using the cyclotomic $L \in \{2, 4\}$ cosine-squared quantisation loss $\mathcal{L}_{\text{pqo}}$ of Remark 3.2 together with the synthetic magnitude-only target Ψ_n of Equation (43).

Outcome. Negative at both $L = 2$ and $L = 4$. At $L = 2$ the network collapses to the trivial attractor $\phi \to \pi/2$ and discriminates neither zeros nor non-zeros. At $L = 4$ the training converges without attractor collapse but the resulting phase MSE against $\arg Z(t)$ is indistinguishable between zeros and non-zeros (pooled Welch $p > 0.12$ across three seeds, cyclotomic sectors $k = 2, 3$ empty). See Remark 3.2 for the full tabulated results.

Diagnosis. We trace the failure through four stages (the dates indicate the diagnostic chain, not the experimental schedule):

1. **Target decoupling (Remark 4.2, 2026-04-08).** The synthetic target Ψ_n of Equation (43) depends only on $\|X\|$ (or on n) and carries no t-local Riemann-zero information. No choice of $(\alpha, \beta, \gamma, \delta)$ repairs this; the defect is structural.

2. **Supervised replacement is necessary but not sufficient (Remark 4.2, 2026-04-08).** Replacing the target by the analytic ground truth $\Psi(t) = \arg \xi(\frac{1}{2} + it)$ and removing $\mathcal{L}_{\text{pqo}}$ from training (keeping it only as measurement) unlocks the asymmetric basin collapse of $L = 4$ but leaves the headline separation metric near the null ($p = 0.976$, ratio 1.004).

3. **Loss-weight imbalance and arithmetic mean read-out are additional broken links (Remark 4.2, 2026-04-08).** The residual loss weights $\lambda_{\text{res}} = 1$, $\lambda_{\text{curv}} = 0.1$ dominate the supervisory $\mathcal{L}_{\text{tgt}}$, and the channel-wise arithmetic mean used to read out $\hat{\Psi}$ destroys the $\pm\pi$ wrap signal predicted by the phase-jump structure on the critical line.

4. **Isolation confirms the chain (Remark 4.3, 2026-04-09).** Setting $\lambda_{\text{res}} = \lambda_{\text{curv}} = \lambda_{\text{pqo}} = 0$ and reading out via a circular mean restores supervised learning with $p = 1.18 \times 10^{-15}$ on the off-zero set, leaving a residual on-zero MSE of $\approx (\pi/2)^2$ explained entirely

by the smooth network's inability to represent the $\pm\pi$ discontinuity of $\arg \xi(\frac{1}{2} + it)$ at each zero.

Lesson. *Four independent failure modes can stack to produce a single negative result.* Before declaring that an analytic identification is experimentally refuted, every link of the supervisory chain must be individually isolated: target injection, loss weighting, measurement operator, and output-space expressivity. The analytic content of the programme $\Psi(t) = \arg \xi(\frac{1}{2} + it)$ survives the cascade; the smooth-field output representation was the last obstruction, *resolved* on 2026-04-10 by replacing the $\{0, \pi\}$ step target with a normalised smooth surrogate $\operatorname{Re} \xi / \max |\operatorname{Re} \xi|$, which reduces the zero/non-zero MSE ratio from 18.2× to 0.03× (Remark 4.3). The complex-vector readout head remains an untested alternative (Section 15.3).

12 ■ Phase-Gauge Gradient Descent (PGGD)

Programme. PGGD is a custom optimiser introduced to exploit the S^1-valued phase structure of the discriminant network's intermediate state. Each parameter update is projected onto the tangent space of the ambient circle bundle so that the optimisation trajectory respects the gauge symmetry $\phi \mapsto \phi + \text{const}$. The hypothesis was that PGGD would outperform Adam on the F_2-based zero detection task of Section 8 by eliminating the spurious global-phase degree of freedom.

Outcome. Negative relative to Adam on every seed tested. PGGD either (i) fails to converge within the same epoch budget, or (ii) converges to a higher training loss with a discriminant gap smaller than Adam's baseline. On the $t \in [10, 50]$ benchmark, Adam's final F_2-MSE is consistently 2–4× smaller than PGGD's best run.

Diagnosis. Two related mechanisms:

1. **Ergodic diffusion on S^1.** The tangent projection used by PGGD removes the radial component of each update but preserves the full angular component. Over many steps the angular update accumulates approximately diffusively, with effective step size η and variance growing linearly in the epoch count. The projected trajectory therefore explores S^1 essentially uniformly (Weyl equidistribution of angles modulo 2π), defeating the very gauge-removal the optimiser was designed to implement. What PGGD gains by removing the radial degree of freedom, it loses to a random walk in the angular one.

2. **Gradient scale mismatch with $\mathcal{L}_{\text{pqo}}$.** The cosine quantisation loss has $O(1)$ gradient away from its minima and $O(|\phi - \phi_*|)$ gradient near them. PGGD's tangent projection amplifies the first regime relative to Adam's adaptive per-parameter scaling, producing an effective learning rate that is too large near the attractors.

Lesson. *Gauge-respecting optimisers do not automatically inherit the sample efficiency of adaptive first-order methods.* A correct tangent projection is compatible with, but does not imply, coherent detection. The PGGD programme is retained as an idea for possible rehabilitation via adaptive per-angle step sizes (see Section 15.3), but in its present form it should not be used in place of Adam on the F_2 detector.

Part IV

Synthesis

13 Structural Correspondence Table

We formalise the seven structural correspondences identified throughout this work:

Table 5: Seven structural correspondences between QROP-Net and the resonance framework.

#	QROP-Net (Concrete)	Resonance (Analytic)	Shared Mathematical Structure
1	$\sin^2(2^{d-1}\phi)$ quantisation loss	$\Delta v = m\pi$ phase quantisation (Gate A)	Periodic phase penalty: $\mathcal{P}(\phi; L) = \sin^2(L\phi/2) = 0$ at lattice positions
2	NTT (discrete) + FrFT (continuous fractional)	Fourier analysis of $\xi(s)$ on the critical line	Spectral decomposition: algebraic structure $\leftrightarrow$ frequency domain
3	$X^n + 1 = 0$ roots: n equally-spaced phases on unit circle	Phase winding $\oint \nabla v \cdot ds = 2\pi N$	Cyclotomic/circular quantisation: discrete phases on S^1
4	LWE error threshold: $\lvert e\rvert_{\mathbb{Z}_q} < q/4 = 1920$	σ-lift: $\lvert\sigma^* - 1/2\rvert = 0$ (RH)	Hard error boundary beyond which signal recovery fails
5	$U_0 = e^{i\phi}$, $\lvert U_0\rvert \equiv 1$; intensity $I = \lvert U_d\rvert^2$	$A(t) = \lvert\xi\rvert$ vs $\Phi_\xi(t) = \arg\xi$	Amplitude–phase separation: information encoded in phase, measured via amplitude
6	Blind TTO: self-calibration without ground truth	Gauge ODE: $F_2 = 0$ without knowing zero locations	Ground-truth-free convergence via internal consistency
7	$\Delta d = 0.1\,\mathrm{mm} \to$ DFR $= 1$	σ-lift $\neq 0 \to$ non-unitarity (absorption)	Sensitivity threshold: small parameter shift $\to$ catastrophic failure

13.1 Mathematical Formalisation of the Correspondences

We express each correspondence as a structural mapping.

Definition 13.1 (Phase Quantisation Functor)**.** Let $(\mathcal{A}, \Phi_\mathcal{A}, L_\mathcal{A})$ be a *phase quantisation triple* consisting of an algebraic structure $\mathcal{A}$, a phase function $\Phi_\mathcal{A} : \mathcal{A} \to [0, 2\pi)$, and a lattice level count $L_\mathcal{A} \in \mathbb{N}$. The phase quantisation operator $\mathcal{P}(\cdot; L)$ acts identically on all such triples:

	$\mathcal{A}$	$\Phi_\mathcal{A}$	$L_\mathcal{A}$
QROP-Net	$\mathbb{Z}_q^{1024}$	$\phi_i = 2\pi c_i/2^{d_u}$	$2^{d_u} = 2048$
Cyclotomic	R_q	$\arg(\omega_j) = \pi(2j+1)/n$	$2n = 512$
Critical line	$\{t \in \mathbb{R}^+\}$	$\vartheta(t) \bmod \pi$	2 (per Gram interval)

In each case, the condition $\mathcal{P}(\Phi_\mathcal{A}; L_\mathcal{A}) = 0$ selects the "aligned" elements, and $\nabla\mathcal{P}$ drives misaligned elements toward the nearest lattice position.

13.2 ■ Dynamical and spectral correspondences

Tier 2. The correspondence table above is structural. This subsection records four additional correspondences that connect the empirical phenomena of Sections 8–9 to classical mathematical frameworks. Each is established at the level of analogy and partial numerical agreement; we do not claim rigorous equivalences.

13.2.1 Maslov correction as the metaplectic sheet shift

The replacement of $\sin^2(L\phi/2)$ by $\cos^2(L\phi/2)$ as the quantisation penalty — the "$\cos^2$ Maslov correction" invoked throughout Section 8 — is the neural shadow of a classical metaplectic phenomenon.

Theorem 13.1 (Metaplectic correction of Gate A)**.** *For the phase quantisation operator* $\mathsf{PQO}(\varphi; L)$ *applied at zeros of a real-valued function where* $\varphi \to \pi/2$ *(Corollary 8.2), the* $\sin^2$ *form gives* $\mathsf{PQO}(\pi/2; 2) = 1$ *(maximal penalty), while the* $\cos^2$ *form gives*

$$\mathsf{PQO}_{\cos}(\varphi; L) := \cos^2(L\varphi/2), \quad \mathsf{PQO}_{\cos}(\pi/2; 2) = 0. \tag{83}$$

The replacement $\sin^2 \to \cos^2$ *is the* Maslov correction *with index* $\mu = 2$*, corresponding to a phase shift of* $\mu\pi/4 = \pi/2$*.*

Table 6: The $\sin^2 \leftrightarrow \cos^2$ duality across mathematics.

Domain	$\sin^2$ **regime**	$\cos^2$ **regime**	**Name of shift**
WKB/semiclassical	Away from turning points	At turning points	Maslov index μ
Symplectic geometry	Standard Lagrangian chart	Crossed Maslov cycle	$\pi_1(\Lambda(n))$
Harmonic analysis	In-phase (I)	Quadrature (Q)	Hilbert transform
Lattice cryptography	Lattice decoding	Coset decoding	$\Lambda \to \Lambda + v$
Berry–Keating	$\vartheta(t_n) = n\pi$	$\vartheta(t_n) = (n - \frac{1}{2})\pi$	half-integer shift
Representation theory	Compact Cartan	Split Cartan	Cayley transform

The unifying object is the metaplectic group $\mathrm{Mp}(2n, \mathbb{R})$, the double cover of $\mathrm{Sp}(2n, \mathbb{R})$, whose two sheets correspond exactly to the two quantisation phases. The functional equation $\xi(s) = \xi(1 - s)$ arises from the action of this group on theta functions via $\tau \to -1/\tau$, and the Berry–Keating quantisation condition

$$\vartheta(t_n) = (n - \tfrac{1}{2})\pi + O(1/t_n) \tag{84}$$

encodes the same $-1/2$ shift on the analytic side.

13.2.2 Gauge ODE critical transition as Kuramoto synchronisation

Observation 13.2 (Gauge critical transition)**.** *In the* F_2*-detector of Section 8, the gauge buffer states* (φ, ψ) *exhibit a sharp transition at* $\sim$10 *epochs: the standard deviation* $\mathrm{std}(\varphi)$ *drops from* ≈ 0.10 *to* ≈ 0.03*, coinciding with the epoch at which* F_2 *reaches convergence on*

all training zeros. The transition is invariant across seeds and has been reproduced on both $t \in [10, 50]$ *and* $t \in [100, 200]$.

Remark 13.1 (Kuramoto synchronisation)**.** Observation 13.2 is a finite-population Kuramoto-type synchronisation transition. The gauge phases $\{\varphi_i\}$ are batch-local fields; the shared residual loss $\mathcal{L}_{\text{res}}$ acts as a coupling functional between them. When the effective coupling strength (controlled by the learning rate and the sharpness of the target) exceeds a critical value, the phases lock onto a common frequency profile determined by the target zero locations. The std(φ) drop of $\approx 3\times$ at epoch ~ 10 is the neural-network signature of this locking.

[confirmed 2026-04-12] The Kuramoto order parameter $r(e) = |N^{-1} \sum_j e^{i\varphi_j}|$ was measured across 3 seeds $\times$ 2 t-ranges (5 complete runs + 1 partial at epoch 90/100 due to OOM; the partial run shows the same monotonic pattern). The std(φ) ratio is $3.43 \pm 0.09\times$ (initial 0.108 ± 0.008, final 0.032 ± 0.002), quantitatively matching the $\sim 3\times$ observation above. The order parameter satisfies $r > 0.99$ from epoch 0 ($r_{\text{init}} = 0.994 \pm 0.001$, $r_{\text{final}} = 0.9994 \pm 0.0001$), indicating that the network initialisation already places the gauge phases in the supercritical regime $K \gg K_c$; training refines the synchronisation rather than triggering a disorder-to-order transition. Both $t \in [10, 50]$ and $t \in [100, 200]$ yield identical patterns, confirming t-range invariance.

Remark 13.2 (von Mises self-consistency)**.** The final phase dispersion $\sigma_\varphi = 0.032 \pm 0.002$ is exactly consistent with a von Mises distribution of concentration parameter $\kappa = 1/\sigma_\varphi^2 \approx 980$: the self-consistency relation

$$r \;=\; e^{-\sigma^2/2} \;=\; e^{-0.032^2/2} \;\approx\; 0.9995$$

matches the measured $r_{\text{final}} = 0.9994 \pm 0.0001$ within experimental uncertainty. This identifies the trained gauge-phase ensemble as a *near-degenerate* von Mises distribution and establishes that the Kuramoto synchronisation is not merely an analogy: the phases are statistically von Mises–distributed with $\kappa \approx 10^3$. [confirmed 2026-04-13]

Proposition 13.3 (Spontaneous gauge fixing)**.** *Let H denote the number of hidden channels. The H gauge phases $\{\varphi_j(x)\}_{j=1}^H$ satisfy $r_{\text{init}} \approx 1 - c\,H^{-\alpha}$ at initialisation, with $\alpha \approx 0.38$ and $c \approx 0.033$ (5-seed fit, $R^2 = 0.94$). Training refines the concentration to $r_{\text{final}} > 0.998$ uniformly, effecting a spontaneous symmetry breaking $U(1)^H \to U(1)$.*

H-scaling data (5 seeds $\times$ 5 widths, $t \in [100, 200]$, 30 epochs):

H	r_{init}	$1-r_{\text{init}}$	r_{final}	$\Delta r/\text{ep}$ ($\times 10^{-3}$)
16	0.9878 ± 0.0016	0.0122	0.9982 ± 0.0003	0.348 ± 0.043
32	0.9915 ± 0.0011	0.0085	0.9985 ± 0.0001	0.236 ± 0.038
64	0.9938 ± 0.0010	0.0062	0.9988 ± 0.0002	0.166 ± 0.027
128	0.9954 ± 0.0003	0.0046	0.9987 ± 0.0001	0.110 ± 0.010
256	0.9955 ± 0.0005	0.0045	0.9989 ± 0.0001	0.112 ± 0.014

The scaling saturates at $H \geq 128$ ($1-r_{\text{init}} \approx 0.0045$), ruling out a simple $O(1/H)$ law. Even at $H = 16$, $r_{\text{init}} = 0.988 \gg 1/\sqrt{16} = 0.25$: the near-unity synchronisation is structural

(driven by Xavier initialisation and ReLU non-linearity suppressing output-phase diversity) rather than a finite-size artifact. The refinement rate $\Delta r/\mathrm{ep}$ *decreases monotonically with* H*, consistent with diminishing marginal gain as* $r_{\mathrm{init}} \to 1$*. [confirmed* $H \in \{16, 32, 64, 128, 256\}$*, 2026-04-13]*

Remark 13.3 (von Mises–Kuramoto–L_{geo} triangle)**.** The geodesic loss $L_{\mathrm{geo}} = 1 - \cos(\Delta\varphi)$ is, up to an additive constant, the negative log-likelihood of the von Mises distribution: $-\log p(\varphi \mid \mu, \kappa) = -\kappa \cos(\varphi - \mu) + \mathrm{const}$. Minimising L_{geo} therefore performs von Mises MLE, which maximises κ, which in turn drives the Kuramoto order parameter $r \to 1$. The three concepts—von Mises MLE, S^1 geodesic loss, and Kuramoto synchronisation—are thus different views of a single optimisation problem. This explains why replacing $\mathcal{L}_{\mathrm{tgt}}$ with L_{geo} (Section 15.3, Tier 1) yields a 4× detection improvement under the original threshold-based metric: the loss function implicitly reinforces gauge synchronisation. A follow-up minima-based evaluation (5 seeds each, two height ranges) reveals a *density-dependent* effect:

- $t \in [500, 600]$ (density 0.72/unit): F_1 comparable (BL 0.373 ± 0.020 vs L_{geo} 0.361 ± 0.012, p=0.45); recall-variance reduction significant (F=16.2, p=0.010).
- $t \in [1000, 1100]$ (density 0.81/unit): F_1 improves (BL 0.395 ± 0.009 vs L_{geo} 0.411 ± 0.005, paired t=2.84, p=0.047, 5/5 seeds directionally consistent); recall-variance also significant (F=6.5, p=0.049).

Across both ranges, L_{geo} stabilises the recall floor ($R \geq 0.972$ at $t \in [500, 600]$; $R \geq 0.975$ at $t \in [1000, 1100]$; all 10 seed–range combinations). At higher zero density the geodesic loss additionally lifts mean F_1, suggesting that the $\pm\pi$-discontinuity burden of $\mathcal{L}_{\mathrm{tgt}}$ grows with density. The effect sizes are small ($\Delta F_1 \approx 0.016$) and the p-values borderline; we report these as confirmed but modest improvements. Catastrophic divergence occurred in 1/5 seeds (seed-specific, mitigated by best-checkpoint selection). [confirmed 2026-04-13, 5-seed, 2 ranges]

13.2.3 Lehmer phenomenon and GUE spacing statistics

The difficulty hierarchy of Table 1, where some zeros converge one to two orders of magnitude slower than their neighbours, matches the empirical distribution of Lehmer pairs. Montgomery's pair-correlation conjecture and the numerical work of Odlyzko establish that the normalised spacings of high Riemann zeros follow the GUE distribution

$$p_{\mathrm{GUE}}(s) = \frac{32}{\pi^2} s^2 e^{-4s^2/\pi}, \tag{85}$$

with close zero pairs — Lehmer pairs — occurring with predictable frequency. In our detector, these are the zeros at which $Z(t)$ nearly grazes the real axis without crossing; the analytic signal spends longer near $\pm\pi/2$, and the bundle curvature profile becomes qualitatively more complex (broader FWHM, multi-peaked). However, the Hadamard partial-sum convergence rate is insensitive to zero spacing ($|r| < 0.03$). This gives a testable prediction: the distribution of per-zero convergence times at fixed epoch budget should match the convolution of the GUE density with the analytic-signal residence time, a quantity that can be computed in closed form. We have not yet performed this fit; it is recorded as an open problem in Section 15.3.

13.2.4 Weyl equidistribution and the PGGD angular walk

Section 12 attributes the PGGD failure to a Weyl-type equidistribution of the angular updates modulo 2π. The underlying fact is the classical theorem that for any irrational α the sequence $\{n\alpha \bmod 1\}$ is equidistributed in $[0, 1)$. In the PGGD context, the angular update per step is approximately $\eta\,\hat{L}(t)$, which for a generic target function is rationally independent over any finite sample; the tangent projection then concentrates the full update into this angular component, producing an equidistributed walk whose empirical distribution converges to the Haar measure on S^1 as the epoch count grows. This is the converse of the locking in Remark 13.1: here the coupling has been removed by construction, and equidistribution wins.

14 ■ Midpoint Curvature Nonlocality

Tier 1. At the midpoint $m = (\gamma_n + \gamma_{n+1})/2$ between consecutive zeros, the Hadamard product ensures that the two nearest-neighbour (NN) contributions to Λ'/Λ cancel exactly: the terms $1/(s - \rho_n)$ and $1/(s - \rho_{n+1})$ sum to zero by symmetry. The residual curvature $\kappa(m) = |\Lambda'/\Lambda(1/2 + im)|^2$ therefore encodes only *non-local* information from distant zeros.

Observation 14.1 (Midpoint curvature encodes the next gap)**.** *For the first 234 consecutive zero pairs of $\zeta(s)$ ($t \in [0, 450]$, dps = 100, analytic $\log \xi'$ formula, cap rate 0.0%), the Spearman rank correlations are:*

Pair	ρ	*p-value*	***Interpretation***
κ_{mid} *vs. gap* $(\gamma_{n+1} - \gamma_n)$	-0.031	0.64	*NN cancellation confirmed*
κ_{mid} *vs.* gap_{next} $(\gamma_{n+2} - \gamma_{n+1})$	-0.654	7.3×10^{-30}	***strong nonlocal signal***
κ_{mid} *vs. local density*	$+0.022$	0.74	*no density effect*

The partial correlation $\rho(\kappa_{\text{mid}}, \text{gap}_{\text{next}} \mid \text{gap}) = -0.666$ remains essentially unchanged, ruling out gap autocorrelation as a confound. [confirmed 2026-04-14]

Observation 14.2 (Universality across Dirichlet L-functions)**.** *The midpoint nonlocality persists across three Dirichlet L-functions (χ_3, χ_4, χ_5), using the analytic formula $\Lambda'/\Lambda = \frac{1}{2}\log(q/\pi) + \frac{1}{2}\psi((s+a)/2) + L'/L(s, \chi)$. All tested L-functions exhibit $|\rho(\kappa_{\text{mid}}, \text{gap}_{\text{next}})| > 0.55$ with $p < 10^{-10}$. [confirmed 2026-04-14]*

Table 7: Midpoint curvature nonlocality: L-function vs. random point processes. "NN cancel" column: $|\rho(a)| < 0.3$ confirms nearest-neighbour cancellation is operative.

Point process	n **pairs**	$\rho(a)$ gap	$\rho(b)$ gap_{next}	**partial** ρ	**Verdict**
$\zeta(s)$	234	-0.031	$\mathbf{-0.654}$	-0.666	✓ nonlocal
$L(s, \chi_3)$	112	-0.045	$\mathbf{-0.567}$	-0.576	✓ nonlocal
$L(s, \chi_4)$	120	-0.016	$\mathbf{-0.608}$	-0.629	✓ nonlocal
$L(s, \chi_5)$	127	-0.013	$\mathbf{-0.607}$	-0.613	✓ nonlocal
GUE (500×500, 20 runs)	20×399	-0.037	-0.008 ± 0.034	-0.020	× absent
Poisson (234 pts, 20 runs)	20×232	-0.290	-0.090 ± 0.045	-0.098	× geometric

Observation 14.3 (GUE/Poisson contrast)**.** *To test whether the nonlocal signal is a generic feature of repulsive point processes, we computed* $R(m)^2 = (\sum_{non\text{-}NN} 1/(m-\lambda_k))^2$ *— the direct analogue of* κ_{mid} *— on GUE eigenvalues (*500×500 *Hermitian matrices, bulk 80%, sliding-window unfolding, 20 runs) and Poisson points (234 uniform points, 20 runs).*

GUE: $\rho(b) = -0.008\pm0.034$ *(t-test* $p = 0.29$*) — statistically indistinguishable from zero despite GUE's spectral rigidity. Poisson:* $\rho(b) = -0.090 \pm 0.045$ *(*$p = 4.5 \times 10^{-8}$*) — statistically significant but* $|\rho|$ *is* $\sim1/7$ *of the L-function value; this residual arises from a purely geometric effect: the dominant non-NN term* $1/(m-\lambda_{i+1}) \approx -2/(\text{gap}+\text{gap}_{\text{next}})$ *couples weakly to* gap_{next}*.*

Conclusion: the midpoint curvature nonlocality is specific to the arithmetic zero placement of the four tested L-functions (ζ*,* χ_3*,* χ_4*,* χ_5*). It is not reproduced by GUE level repulsion or Poisson independence. The mechanism remains unidentified; we record this as an open problem. [confirmed 2026-04-14, 20-run ensemble]*

Observation 14.4 (Nonlocal reach decay and forward–backward sign reversal)**.** *To quantify how far the midpoint curvature "sees," we extend Observation 14.1 by measuring* $\rho(\kappa_{\text{mid}}, \text{gap}_{n+k})$ *for* $k = 0, 1, \ldots, 10$ *(forward) and* $\rho(\kappa_{\text{mid}}, \text{gap}_{n-k})$ *for* $k = 1, \ldots, 10$ *(backward), across all four L-functions.*

Forward decay (*ζ*,* $n = 234$*):

k	ρ_{fwd}	*p-value*	*partial* ρ
0	$-$ 0.031	6.3×10^{-1}	—
1	$-$ **0.658**	$\mathbf{3.7 \times 10^{-30}}$	**−0.668**
2	+ 0.068	3.0×10^{-1}	+0.071
3	+ 0.114	8.6×10^{-2}	+0.121
4	+ 0.057	3.9×10^{-1}	+0.065
5	+ 0.003	9.7×10^{-1}	+0.009

The signal drops by an order of magnitude from k=1 *(*$|\rho| = 0.658$*) to* k=2 *(*$|\rho| = 0.068$*), confirming that the nonlocal coupling is sharply concentrated at the nearest non-adjacent gap.*

Forward–backward sign reversal at k=1 ***(all four L-functions):***

L-function	$\rho_{\text{fwd}}(k{=}1)$	$\rho_{\text{bwd}}(k{=}1)$	*sign flip?*
$\zeta(s)$	−0.658	+0.576	✓
$L(s,\chi_3)$	−0.555	+0.504	✓
$L(s,\chi_4)$	−0.608	+0.588	✓
$L(s,\chi_5)$	−0.606	+0.602	✓

For every tested L-function, the forward correlation at k=1 *is* negative *while the backward correlation is* positive*, with comparable magnitudes (*$||\rho_{\text{fwd}}|-|\rho_{\text{bwd}}|| < 0.08$*). This sign reversal is a direct consequence of the Hadamard structure: from the midpoint* $m = (\gamma_n + \gamma_{n+1})/2$*, the dominant non-NN contribution on the right (*γ_{n+2}*) and on the left (*γ_{n-1}*) enter* Λ'/Λ *with opposite signs.*

Decay model. *For the* $|\rho_{\text{fwd}}|$ *sequence (*$k = 1, \ldots, 10$*), power-law fits* $|\rho| \sim A\,k^{-\alpha}$ *consistently outperform exponential fits (*R^2*:* ζ *0.71,* χ_4 *0.76,* χ_5 *0.76;* χ_3 *is a 3-seed preliminary with*

$R^2 = 0.38$). The rapid decay means that a single characteristic length ξ is difficult to define; the nonlocal signal is concentrated almost entirely at k=1. [confirmed 2026-04-14, 4 L-functions, 3-seed preliminary]

15 Discussion

15.1 What the Correspondence Means

The seven correspondences demonstrate a *structural isomorphism* at the level of mathematical grammar: the same pattern of periodic phase penalty, hard error boundary, amplitude–phase separation, and self-calibrating convergence appears in both the concrete engineering system (QROP-Net) and the abstract analytic framework (resonance theory on the critical line).

This suggests—but does **not prove**—a deeper connection between:

1. The algebraic structure of cyclotomic polynomial rings ($X^n + 1 = 0$, the foundation of lattice cryptography),
2. The spectral structure of the Riemann ξ-function (the encoding of prime distribution in complex-analytic phase),
3. The physical structure of coherent optical propagation (the Fresnel diffraction kernel as a fractional Fourier transform).

The unifying thread is the *unit circle*: polynomial roots on S^1, optical phases on S^1, and ξ-function values on S^1 (for the phase component). Phase quantisation—forcing continuous signals onto discrete angular positions—is the shared operation.

15.2 What the Correspondence Does Not Mean

We must be explicit about the boundaries of this work:

1. **This is not a proof of the Riemann Hypothesis.** The structural parallels identify a shared mathematical grammar, but they do not constitute a logical proof. The RH is a statement about the zeros of a specific analytic function, and no amount of structural analogy to an optical system can replace a rigorous complex-analytic argument.
2. **The correspondences are structural, not causal.** The fact that QROP-Net's $\sin^2$ loss and Gate A's $\Delta v = m\pi$ have the same functional form does not mean that one "causes" or "explains" the other. They are independent instantiations of the same mathematical pattern.
3. **The polyadic frame is a numerical tool, not a proof technique.** The convergence of $F_2(t) = 0$ to known zero locations validates the numerical method, but finding zeros is not the same as proving they all lie on $\sigma = 1/2$.
4. **The scattering interpretation is an analogy.** Treating $\xi(s)$ as a scattering amplitude is a productive analogy that generates correct predictions (zero locations, phase quantisation), but ξ is not literally a scattering amplitude of a physical system (no Hamiltonian has been identified whose S-matrix equals ξ).

15.3 Predictive Power and Future Directions

Despite these caveats, the unified framework suggests several testable predictions. We organise them by tier, preserving the classification principle of Section 1: items are grouped by how concrete the evidence already is.

■ **Tier 1 follow-ups (immediate).**

- **Complex-vector readout for** $\arg \xi(\frac{1}{2} + it)$ **[neutral 2026-04-12].** The isolation experiment (Remark 4.3) localised the remaining defect to the smooth network's inability to represent the $\pm\pi$ discontinuity of $\arg \xi$ at zeros. The cleanest architectural fix is to have the readout head predict $\xi(\frac{1}{2} + it)$ as a complex vector and compute arg post-hoc, rather than predict the phase directly. *Tested 2026-04-12:* a linear readout head mapping pre-projection hidden states $Z \in \mathbb{C}^{64}$ to $(\operatorname{Re} \xi, \operatorname{Im} \xi)$ was trained jointly with auxiliary resonance losses ($\mathcal{L}_{\text{res}} + \mathcal{L}_{\text{curv}} + \mathcal{L}_{\text{pqo}}$, $\mathcal{L}_{\text{tgt}} = 0$), using MSE on the complex components ($\lambda_{\text{MSE}} = 2$). Over 3 seeds ($t \in [100, 200]$, $K = 128$, 150 epochs): complex-vector validation loss improved $\sim 4\times$ (0.014 ± 0.005 vs 0.069 ± 0.013), but the $|F_2|$ ratio was statistically unchanged (0.982 ± 0.041 vs 0.999 ± 0.005; $p > 0.5$) and zero detection remained at 0/463. The readout head learns the Re/Im mapping well but does not alter the network's internal resonance structure (Z_{out}). The $\pm\pi$ discontinuity thus requires a deeper architectural change than a post-hoc readout head alone.

- **Smooth zero-crossing targets [confirmed 2026-04-10].** Two smooth surrogates were tested under the isolation setup of Remark 4.3: a bump target $\operatorname{sign}(\operatorname{Re} \xi) \cdot e^{-(\operatorname{Re} \xi)^2/\sigma^2}$ and a normalised target $\operatorname{Re} \xi / \max |\operatorname{Re} \xi|$. The normalised target reduces the zero/non-zero MSE ratio from $18.2\times$ (step target) to $0.03\times$ — a 99.8% gap reduction — confirming that the $\pm\pi$ discontinuity was the sole bottleneck (Remark 4.3). *Full-pipeline test [completed 2026-04-10]:* integrating the smooth target into the full loss ($\mathcal{L}_{\text{res}} + \mathcal{L}_{\text{curv}} + \mathcal{L}_{\text{pqo}}$ + MSE smooth) *worsened* global $|F_2|$ by 5.2% (3-seed ensemble), indicating that the angular target $\mathcal{L}_{\text{tgt}}$ synergises with $\mathcal{L}_{\text{res}}$ and $\mathcal{L}_{\text{pqo}}$ and should not be replaced wholesale. The smooth target remains valid as an *isolation diagnostic* but not as a drop-in substitute for the full pipeline.

- **Blind zero extrapolation [confirmed 2026-04-10].** Training on $t \in [100, 150]$ and predicting local $|F_2|$ minima on the unseen interval $[150, 200]$ (and vice versa) yields recall 96−99% across 5 seeds and both $K \in \{64, 128\}$ settings; reverse direction with $K = 128$ achieves $99.1 \pm 1.7\%$ recall. Precision is low ($\sim$15−18%) due to spurious grid minima, but the near-perfect recall demonstrates that the F_2 residual generalises to unseen height ranges and is not merely fitting training-set zeros. *Post-hoc precision filtering [neutral 2026-04-12]:* ensemble consensus and depth/spacing filters were applied to the blind-extrapolation minima ($t \in [150, 200]$, 5 seeds, $K = 128$, 200 epochs, 2000-point dense grid). The best filter (ensemble $\geq$ 2 seeds) achieves precision 32.7%, recall 63.0%, $F_1 = 0.430$. Precision rises from $\sim$15% (raw) to 28−33% but does not reach the $1.5\times$ significance threshold. Depth filtering is too aggressive (removes all candidates), while spacing filters reduce recall without compensating precision. The low precision appears intrinsic to the $|F_2|$ landscape rather than correctable by post-hoc

heuristics. *FP anatomy [confirmed 2026-04-13]:* Over 5 seeds ($t \in [14, 50]$), 85.1% of false positives (57/67) are genuine local minima of the $|F_2|$ landscape ($|F_2|''(t) > 0$, $p < 10^{-6}$), confirming that pseudo-zeros are structural features of the learned residual, not noise. FP positions are seed-inconsistent (Jaccard = 0.046, $\geq$5-seed consensus = 0), indicating that the fine structure of $|F_2|$ depends on model initialisation. The analytic curvature $\kappa = |\xi'/\xi|^2$ separates TP from FP by a factor of $\sim$300$\times$ (TP median 270, FP median 0.9; gap: FP max = 31 < TP min = 76), establishing a theoretical precision ceiling of 100% if κ were available. Since ξ is real on $\sigma = 1/2$, the monodromy at zeros is trivially $\pm\pi$ (sign change), contributing no new information beyond the curvature separation.

- S^1 **geodesic phase prediction [confirmed 2026-04-12].** Instead of predicting $\arg \xi$ as a scalar (which suffers a $\pm\pi$ branch cut at zeros), the readout head predicts $(\cos\theta, \sin\theta) \in S^1$ and is trained with a geodesic loss $L_{\text{geo}} = 1 - \cos(\theta_{\text{pred}} - \theta_{\text{true}})$. Over 3 seeds ($t \in [100, 200]$, $K = 128$, 150 epochs): S^1 detection improves from 5.7 to 13.7 per 463 zeros (2.4$\times$); angular error at zeros is 0.215 rad vs 0.155 at non-zeros (1.39$\times$ ratio), confirming that the $\pm\pi$ discontinuity is partially bypassed. The $|F_2|$ ratio is unchanged (1.054 ± 0.103 vs 1.026 ± 0.233); the detection gain comes from the readout head's improved phase tracking near zeros rather than from changes to the network's resonance structure. *Follow-up (completed):* Replacing $\mathcal{L}_{\text{tgt}}$ entirely with L_{geo} in the full loss ($\mathcal{L}_{\text{res}} + \mathcal{L}_{\text{curv}} + L_{\text{geo}} + \mathcal{L}_{\text{pqo}}$) yields a further 1.66$\times$ detection gain over the readout-only approach: 22.7/463 zeros detected vs 13.7 (readout) vs 5.7 (baseline), a cumulative 4$\times$ improvement. The $|F_2|$ ratio drops to 0.913 ± 0.105 (from 1.026 ± 0.233), meaning zeros now exhibit *smaller* $|F_2|$ residuals than non-zeros—a qualitative inversion of the baseline regime. Seed-level variance is high (5, 40, 23 detections), suggesting that the geodesic loss landscape has multiple basins with differing phase-alignment quality. [confirmed 2026-04-12, threshold metric] *Follow-up (completed, 5 seeds):* A minima-based evaluation at $t \in [500, 600]$ with 5 seeds $(42, 7, 123, 314, 2024)$ shows that the 4$\times$ improvement is metric-conditional. Under local-minimum detection, F_1 is comparable (BL 0.373 ± 0.020 vs L_{geo} 0.361 ± 0.012; difference non-significant). The only consistent effect of L_{geo} is **recall-floor stabilisation**: BL recall 0.956 ± 0.045 (range $[0.889, 1.000]$) while L_{geo} maintains 0.978 ± 0.011 (range $[0.972, 1.000]$), a 4.1$\times$ variance reduction. An F-test on the variance ratio yields $F = 16.7$ (df $= (4, 4)$, $p < 0.05$), confirming statistical significance. The worst-case seed 2024 shows $R_{\text{BL}} = 0.889 \rightarrow R_{L_{\text{geo}}} = 0.972$ (+3 zeros out of 36), demonstrating that L_{geo} provides a *recall floor* rather than a mean improvement. The $|F_2|$ ratio improvement (0.913 vs 1.026) persists, but does not translate into F_1 gains under minima-based evaluation. [confirmed 2026-04-13, 5-seed F-test significant] *Extension to* $t \in$ [1000,1100] *[confirmed 2026-04-13, 5-seed, recall F-test + F_1 paired t-test significant]:* The interval contains 81 zeros (density 0.81/unit vs 0.72/unit at $t \in [500, 600]$), split into $[1000, 1050]$ (training, 41 zeros) and $[1050, 1100]$ (blind prediction, 40 zeros).

  Baseline (5 seeds): R=0.965 ± 0.029, F_1=0.395 ± 0.009. Val-loss instability is systematic: 3/5 seeds exhibit ep125 spikes > 1.0. Notably, seed 2024 achieves R=1.000—showing

that BL can also reach perfect recall, refuting the claim that R=1.000 is unique to L_{geo}.

L_{geo} *(5 seeds):* R=0.995±0.011, F_1=0.411±0.005. Catastrophic divergence in 1/5 seeds (seed 42, val= 306.93; mitigated by best-checkpoint selection, which preserves the same R=1.000 and F_1=0.412). Remaining 4 seeds stable or mildly divergent.

Statistical tests: F_1: paired t=2.84, p=0.047; all 5 seeds directionally consistent ($\Delta F_1 \in [+0.006, +0.036]$, sign test p=0.031). **Recall variance**: F=6.5, p=0.049 (one-tailed), consistent with the F=16.2 (p=0.010) at $t \in [500, 600]$. $R \geq 0.975$ **floor**: maintained across all 5 seeds. The density-dependent pattern—recall-variance reduction at both densities, but F_1 improvement only at the higher density (0.81/unit)—suggests that the $\pm\pi$ discontinuity burden of $\mathcal{L}_{\text{tgt}}$ grows with zero density. Effect sizes are small ($\Delta F_1 \approx 0.016$) and p-values borderline; we report these as confirmed but modest improvements.

- **Tier 2 programmes (conjectural but well-posed).**

1. **Height scaling law for $\overline{|F_2|}$ [confirmed 2026-04-12].** Blind zero extrapolation was tested at three height ranges: $t \in [100, 200]$ (baseline), $t \in [500, 600]$, and $t \in [1000, 1100]$, each split into a training half and a blind prediction half ($K \in \{128, 256\}$, 3 seeds, 150 epochs, 2000-point dense grid). Recall is stable across a 10× height increase: 98.8±1.7% at t∼200, 98.1 ± 1.3% at t∼600, 97.5 ± 2.0% at t∼1100. Precision *increases* with height (17.1% → 22.8% → 25.6%), likely because the rising zero density reduces the fraction of spurious minima. $K = 256$ offers no advantage over $K = 128$ at any height. The F_2 residual thus scales to heights an order of magnitude beyond the original training range without requiring bandwidth adjustment.

2. **Architecture-invariant residual.** Define a bootstrap-minimum version of $|F_2(t_k)|$ across a model ensemble (Section 10.4) and test whether its per-zero values are now reproducible.

3. **Kuramoto order parameter.** [completed 2026-04-12] The Kuramoto order parameter r was measured over 3 seeds × 2 t-ranges (5 complete + 1 partial). Key findings: std(φ) ratio 3.43±0.09× (quantitative match to Observation 13.2), $r_{\text{init}} = 0.994$, $r_{\text{final}} = 0.9994$. The network starts in the supercritical regime ($r > 0.99$ at epoch 0); training is phase refinement, not a disorder-to-order transition. See Remark 13.1.

4. **Lehmer/GUE convergence fit.** Test whether the per-zero convergence rate matches the convolution of the GUE spacing density with the analytic-signal residence time, as predicted in Section 13.2.3.

5. **Generalised resonance on automorphic L-functions.** The polyadic frame extends naturally to Dirichlet and modular L-functions; the gate conditions should carry over with no modification beyond replacing the Riemann–Siegel ϑ by the appropriate theta function. [confirmed 2026-04-13: bundle properties verified for χ mod 3, 4, 5; see Appendix C.1.] *Blind zero prediction [confirmed 2026-04-13]:* Curvature-peak and monodromy-based detection on the unseen interval $t \in [25, 40]$ achieves F1 =

1.000 for all three characters (χ mod 3, χ mod 4, χ mod 5): 21/21 zeros detected with tolerance $\delta t < 0.5$, every predicted point satisfying $|\text{mono}|/\pi = 1.0000$. See Appendix C.2. *Cross-comparison (partial):* $\sigma = 0.5$ uniqueness (phase jumps and energy concentration) confirmed for all four functions (ζ and three Dirichlet L-functions); curvature/monodromy comparison pending zero-finding fix.

6. **FP monodromy anatomy.** The FP anatomy experiment (Section 15.3, Tier 1) established that κ separates TP from FP by $\sim$300×. A pending rerun with contour-integral monodromy (replacing the earlier eps-differentiation method) will test whether $|\text{mono}|/\pi \approx 0$ at all FP locations, which would establish a clean binary classifier: monodromy $= \pm\pi \Rightarrow$ true zero, monodromy $\approx 0 \Rightarrow$ false positive. [pending 2026-04-13]

7. **Higher-conductor Dirichlet extension.** Extension to $q \in \{7, 8, 11\}$ confirms bundle universality beyond small conductors. All three new characters (χ_7 complex order 6, χ_8 real, χ_{11} complex order 10) satisfy: curvature concentration κ-ratio 395−654×, monodromy quantization $|\text{mono}|/\pi = 2.0000$ (contour integral; simple zero $\Rightarrow$ winding number 1), energy ratio $E(0.5)/E(0.3) = 194-331\times$. Phase-jump uniqueness is *fully confirmed*: a dedicated off-critical anatomy (branch-cut–corrected arg with threshold $\pi/2$, matching the ζ protocol) yields $\max_{\sigma\neq 0.5} = 0$ for all three characters (Λ and bare L alike). The earlier apparent weakening ($\max_{\sigma\neq 0.5} = 8$) was a measurement artefact caused by uncorrected arg branch-cut wrapping in the preliminary script; all five bundle properties now hold at higher conductors. The real-character energy anomaly is confirmed general: χ_4: 132.5×, χ_8: 194.3× (both self-dual), vs complex characters 12−331×. [confirmed 2026-04-13]

■ **Tier 3 rehabilitations (failed programmes that could be revisited).**

- **PGGD with adaptive per-angle step sizes.** Section 12 identifies Weyl equidistribution as the cause of PGGD's failure. A per-angle adaptive learning rate that shrinks with the accumulated angular drift would directly counter the diffusion, potentially rehabilitating the optimiser.

- **Cyclotomic PQO with supervised Ψ.** The $L = 4$ programme of Section 11 became statistically significant only after all three defects of Remark 4.2 were closed. A clean rerun of the original cyclotomic framework with the isolation-style setup would produce a fair test of the structural prediction rather than of the now-understood implementation bugs.

QROP-Net and cross-domain predictions (unchanged from the original framework).

1. **Loss landscape transfer**: Can QROP-Net-style training losses (fidelity + $\sin^2$ quantisation + exponential barrier) improve numerical zero-finding algorithms for $\zeta(s)$?

2. **Deep learning on ξ**: Can a neural network trained on $|\xi(\sigma + it)|^2$ (analogous to the diffraction intensity) recover $\arg \xi$ (analogous to phase retrieval)?

3. **Cyclotomic primes**: The NTT-friendly prime $q = 7681 \equiv 1 \pmod{512}$ was chosen for computational efficiency. Is there a deeper number-theoretic reason why primes of this form are especially suitable for phase quantisation?

4. **Generalised resonance**: The polyadic frame extends naturally to Dirichlet L-functions. Can the gate conditions be applied to automorphic L-functions, where the generalised Riemann Hypothesis (GRH) is also conjectured?

5. **Physical implementation**: Can an actual optical experiment with an SLM, laser, and camera sensor demonstrate the $\sin^2$ quantisation landscape *in hardware*, providing a physical analogue computer for phase quantisation problems?

16 Related Work

ML for Riemann zeros. Shanker [23] trained a two-layer neural network on Gram-point features to predict the distance from a Gram point to the next zero, achieving $R^2 > 0.99$ on 80,000 observations. Avysogorets [24] extended this with SVR and RNN regressors over 50 hand-engineered features. Both approaches are direct regression on zero locations; our method differs in using a physics-informed gauge ODE whose residual F_2 vanishes at zeros, providing an interpretable diagnostic channel rather than a black-box predictor.

Fourier feature networks and spectral bias. Tancik et al. [25] showed that mapping inputs through random Fourier features $\gamma(x) = [\sin(2\pi Bx), \cos(2\pi Bx)]$ with $B \sim \mathcal{N}(0, \sigma^2)$ overcomes the spectral bias of standard MLPs [21]. Our Fourier embedding (Section 10.3) uses deterministic equally-spaced frequencies $\omega_k = 2\pi k/W$ instead of random ones. Proposition 10.1 can be viewed as a *deterministic* analogue of their bandwidth-tuning insight: where they sweep σ, we derive the minimum number of frequencies K from the Riemann–von Mangoldt density. A direct comparison [completed 2026-04-11] on $t \in [100, 200]$ with $K = 128$ and $\sigma \in \{1, 10, 100\}$ (3-seed ensemble) shows that deterministic spacing achieves $2{-}2.5\times$ lower validation MSE (8.2×10^{-5} vs. $1.7{-}2.1 \times 10^{-4}$), while random features yield a more uniform zero/non-zero MSE ratio (0.81 vs. 1.62). The deterministic basis is retained as the default.

Equivariant networks. The gauge-equivariant architecture we use is related to the steerable CNN framework of Cohen et al. [26] and its implementation in the `escnn` library [27], which formalises SO(2) equivariance via principal fibre bundles and steerable kernels. Our ResonantBlock implements a U(1) gauge structure via Lie-exponential phase propagation; a future direction is to replace it with `escnn`'s steerable convolution backend for automatic equivariance verification.

Symbolic mathematics with deep learning. Lample and Charton [28] demonstrated that sequence-to-sequence transformers can solve symbolic integration and ODE problems. Their approach is complementary to ours: where they learn symbolic manipulation, our gauge ODE encodes the target equation's structure directly in the architecture.

17 Conclusion

We have presented a unified framework that identifies phase quantisation on algebraic structures as the shared mathematical grammar connecting three domains: lattice-based post-quantum cryptography, coherent optical propagation, and the analytic structure of the Riemann ξ-function.

The concrete nucleus is QROP-Net, a fully differentiable pipeline whose $\sin^2$ quantisation loss, LWE error threshold, and blind self-calibration mechanism find structural counterparts in the two-gate resonance criterion ($\Delta v = m\pi$, $v_\sigma = 0$), the σ-lift from the critical line, and the polyadic frame's gauge ODE for ground-truth-free zero detection.

The project name $k^n = -1$ encodes both the mathematical foundation ($X^{256} \equiv -1$ in Kyber's cyclotomic ring) and the deeper observation that the same cyclotomic structure—discrete phases on the unit circle, related by roots of unity—underlies the NTT in cryptography, the fractional Fourier transform in optics, and the Gram point lattice on the critical line.

Seven structural correspondences have been formalised, and their boundaries clearly delineated: they are analogies grounded in shared mathematical structure, not proofs. But analogies of this precision and breadth are rare, and they suggest that the phase quantisation principle may be a more fundamental organising concept than its individual manifestations in cryptography, optics, or number theory.

References

[1] R. Avanzi *et al.*, "CRYSTALS-Kyber: Algorithm specifications and supporting documentation (version 3.02)," NIST PQC Submission, 2021.

[2] H. M. Ozaktas, Z. Zalevsky, and M. A. Kutay, *The Fractional Fourier Transform with Applications in Optics and Signal Processing.* Wiley, 2001.

[3] Y. Wang, X. Cao, and Q. Zhou, "Uformer: A general U-shaped transformer for image restoration," in *Proc. CVPR*, 2022.

[4] Y. Bengio, N. Léonard, and A. Courville, "Estimating or propagating gradients through stochastic neurons," *arXiv:1308.3432*, 2013.

[5] National Institute of Standards and Technology, "Module-Lattice-Based Key-Encapsulation Mechanism Standard (FIPS 203)," 2024.

[6] J. R. Fienup, "Phase retrieval algorithms: A comparison," *Applied Optics*, vol. 21, no. 15, pp. 2758–2769, 1982.

[7] J. W. Goodman, *Introduction to Fourier Optics*, 3rd ed. Roberts & Company, 2005.

[8] A. D. Wyner, "The wire-tap channel," *Bell System Technical Journal*, vol. 54, no. 8, pp. 1355–1387, 1975.

[9] M. V. Berry, "Quantal phase factors accompanying adiabatic changes," *Proc. Royal Society A*, vol. 392, pp. 45–57, 1984.

[10] E. C. Titchmarsh, *The Theory of the Riemann Zeta-Function*, 2nd ed. (revised by D. R. Heath-Brown). Oxford University Press, 1986.

[11] H. M. Edwards, *Riemann's Zeta Function*. Academic Press, 1974.

[12] A. Khare, "The phase of the Riemann zeta function," *Pramana – J. Phys.*, vol. 48, pp. 537–553, 1997.

[13] A. LeClair and G. Mussardo, "Riemann zeros as quantized energies of a scattering system," *arXiv:2307.01254*, 2023.

[14] A. LeClair, "Analogous results for Generalized and Grand Riemann Hypotheses," *arXiv*, 2024.

[15] Z. Pei, "Spectral flow for the Riemann zeros," *arXiv:2403.19118*, 2025.

[16] J. J. G. Moreta, "Riemann zeros and an exponential potential," 2014.

[17] A. Ivić, *The Riemann Zeta-Function: Theory and Applications*. Dover, 2003.

[18] H. L. Montgomery, "The pair correlation of zeros of the zeta function," *Proc. Symp. Pure Math.*, vol. 24, pp. 181–193, 1973.

[19] J. B. Conrey, "The Riemann Hypothesis," *Notices of the AMS*, vol. 50, no. 3, pp. 341–353, 2003.

[20] A. Jacot, F. Gabriel, and C. Hongler, "Neural tangent kernel: Convergence and generalization in neural networks," in *Advances in Neural Information Processing Systems (NeurIPS)*, 2018.

[21] N. Rahaman, A. Baratin, D. Arpit, F. Draxler, M. Lin, F. Hamprecht, Y. Bengio, and A. Courville, "On the spectral bias of neural networks," in *Proc. ICML*, 2019.

[22] R. Basri, M. Galun, A. Geifman, D. Jacobs, Y. Kasten, and S. Kritchman, "Frequency bias in neural networks for input of non-uniform density," in *Proc. ICML*, 2020.

[23] O. Shanker, "Neural network prediction of Riemann zeta zeros," *Advanced Modeling and Optimization*, vol. 17, no. 1, pp. 45–54, 2015.

[24] M. Avysogorets, J. Bhutani, and P. Beineke, "Machine learning for predicting Riemann zeta zeros," *arXiv:2104.12285*, 2021.

[25] M. Tancik, P. P. Srinivasan, B. Mildenhall, S. Fridovich-Keil, N. Ramasinghe, U. Kamber, and R. Ng, "Fourier features let networks learn high frequency functions in low dimensional domains," in *Advances in Neural Information Processing Systems (NeurIPS)*, 2020.

[26] T. S. Cohen, M. Weiler, B. Kicanaoglu, and M. Welling, "Gauge equivariant convolutional networks and the icosahedral CNN," in *Proc. ICML*, 2019.

[27] G. Cesa, L. Lang, and M. Weiler, "A program to build E(N)-equivariant steerable CNNs," in *Proc. ICLR*, 2022.

[28] G. Lample and F. Charton, "Deep learning for symbolic mathematics," in *Proc. ICLR*, 2020.

A QROP-Net Full Equation Reference

For convenience, we collect the complete set of QROP-Net equations used in the structural correspondences:

1. **Ring**: $R_q = \mathbb{Z}_q[X]/(X^n + 1)$, $n = 256$, $q = 7681$, $X^n \equiv -1$
2. **Encryption**: $\mathbf{u} = \mathbf{A}^\top \mathbf{r} + \mathbf{e}_1$, $v = \mathbf{t}^\top \mathbf{r} + e_2 + \lceil q/2 \rfloor\, m$
3. **Compression**: $\mathrm{Compress}_q(x, d) = \left\lceil 2^d x/q \right\rfloor \bmod 2^d$
4. **Phase modulation**: $\phi_i = 2\pi c_i / 2^{d_u}$, $U_0 = e^{i\phi}$
5. **FrFT**: $U_d = A_\alpha^2 \cdot \mathcal{F}^{-1}[\mathcal{F}[U_0 \cdot C_1] \cdot \mathcal{F}[C_2]] \cdot \Delta x^2$
6. **Quantisation loss**: $\mathcal{L}_{\text{quant}} = \frac{1}{2}[\frac{1}{kn} \sum \sin^2(2^{d_u - 1}\phi_i^{(u)}) + \frac{1}{n} \sum \sin^2(2^{d_v - 1}\phi_j^{(v)})]$
7. **LWE threshold**: $\|w - \lceil q/2 \rfloor\, \hat{m}\|_\infty < 1920$
8. **Blind loss**: $\mathcal{L}_{\text{blind}} = \mathcal{L}_{\text{fid}} + 0.5\mathcal{L}_{\text{quant}} + 0.05\mathcal{L}_{\text{TV}}$

B Numerical Tables for Resonance Gates

Representative zero-finding results from the polyadic frame, across different target functions and parameter settings:

Table 8: Polyadic frame results: invariance under step size sweep.

Target	**Band**	W	N	h sweep	t^* (invariant)
ξ	$T \approx 141.12$	0.002	401	2e−7 to 1e−5	141.123703
ξ	$T \approx 141.12$	0.001	301	2e−7 to 1e−5	141.123713
$L(\chi_3^{\text{odd}})$	$T \approx 140.61$	0.002	61	5e−7	140.611723
$L(\chi_8^{\text{even}})$	$T \approx 141.08$	0.003	81	5e−7	141.083569
$L(\chi_5^{\text{cmplx}}, \chi(2) = i)$	$T \approx 140.97$	0.002	61	2.2e−7	140.966280

The frame invariance ratios across all tests: $r_{\text{frame}} < 10^{-3}$, $r_{\text{channel}} < 0.1$, $r_{\text{window}} < 0.2$. Conjugate pair symmetry for complex characters: $|t_\chi + t_{\bar{\chi}}|/t_\chi < 0.002$.

C Generalisation to Dirichlet L-Functions

The gate conditions and polyadic frame extend naturally to completed Dirichlet L-functions $\Lambda(s, \chi)$. The key adaptations:

1. The completed function $\Lambda(s,\chi) = (q/\pi)^{(s+\epsilon)/2}\Gamma((s+\epsilon)/2)\,L(s,\chi)$ replaces $\xi(s)$.
2. For non-principal characters, $\Lambda(s,\chi)$ is entire (no pole at $s=1$), avoiding the truncation issue of the Euler product.
3. The functional equation $\Lambda(s,\chi) = \varepsilon(\chi)\,\Lambda(1-s,\bar{\chi})$ has the same structural form as $\xi(s) = \xi(1-s)$.
4. Gate A becomes $\Delta v_\chi = m\pi$ where $v_\chi = \arg\Lambda(\sigma+it,\chi)$.
5. Gate B remains $\partial_\sigma v_\chi = 0$.
6. The Generalised Riemann Hypothesis (GRH) is equivalent to $\sigma^* = 1/2$ for all resonant points of $\Lambda(s,\chi)$, for all χ.

LeClair [14] has confirmed that the Bethe-ansatz/phase-quantisation approach extends to "all L-functions based on non-principal Dirichlet characters [and] those based on cusp modular forms," with the only difference being the handling of the pole for the principal character.

C.1 Numerical Verification of Bundle Properties

[completed 2026-04-13] We verify that the four ξ-bundle properties established for $\zeta(s)$ carry over to Dirichlet L-functions, using χ mod 3 (simplest non-trivial), χ mod 4 (real, Kronecker symbol), and χ mod 5 (complex primitive character). All computations use `mpmath` at 80 decimal digits; zeros are located by sign-change bisection of $\Re\Lambda$ and $\Im\Lambda$ on the critical line, refined by `findroot`.

Character	**Zeros**	**Phase-jump isolation**	κ **ratio**	**Mono. dev.**	$E(0.5)/E(0.3)$
χ mod 3 (real)	12	σ=0.5 only	16,457×	0.0000	21.6×
χ mod 4 (real)	13	σ=0.5 only	22,436×	0.0000	454.6×
χ mod 5 (complex)	13	σ=0.5 only	15,732×	0.0000	14.5×

All 38 zeros across three characters exhibit exact monodromy $\Delta\arg\Lambda = \pm\pi$ (deviation $< 10^{-4}$). Phase jumps occur exclusively at $\sigma = 1/2$ ($t \in [10,40]$, 11 values of σ tested). Curvature concentration ratios are computed as median κ at zero-proximal points (± 0.01) vs. generic points (> 1.0 from any zero).

Remarks.

1. The $\pm\pi$ monodromy for the *complex* character χ mod 5 is not a consequence of Λ being real-valued on the critical line (it is not). Rather, it follows from the topology of simple zeros: near a simple zero s_0, $\Lambda(s_0+i\varepsilon)/\Lambda(s_0-i\varepsilon) \to \Lambda'(s_0)\cdot i\varepsilon/[\Lambda'(s_0)\cdot(-i\varepsilon)] = -1$, so $\Delta\arg = \pm\pi$. This is universal for all analytic functions with simple zeros.
2. The energy ratio $E(0.5)/E(0.3)$ varies by a factor of ~ 30 across characters (14.5 to 454.6). The cause is not yet understood and is left as an open observation.

3. These results confirm that the ξ-bundle framework is not specific to $\zeta(s)$ but applies uniformly to the family of Dirichlet L-functions. Each individual property is a consequence of standard analytic number theory (argument principle, GRH); the contribution is the *unified geometric language* that organises them into a single fibre-bundle picture.

C.2 Blind Zero Prediction

[completed 2026-04-13] Using the curvature profile $\kappa(1/2+it)$, monodromy jumps $|\Delta \arg \Lambda| > \pi/2$, and their conjunction, we predict zeros in the test interval $t \in [25, 40]$ for each character.

Character	**True**	**Method**	N_{pred}	**P**	**R**	**F1**
χ mod 3	7	curvature / dual / mono	7 / 7 / 7	1.000	1.000	1.000
χ mod 4	7	curvature / dual / mono	7 / 7 / 7	1.000	1.000	1.000
χ mod 5	7	curvature / dual / mono	7 / 7 / 7	1.000	1.000	1.000

All three methods achieve F1 = 1.000 on all three characters (21/21 zeros, tolerance $\delta t < 0.5$). At 80-digit `mpmath` precision, curvature peaks alone suffice to locate every zero without false positives; the dual and monodromy-only criteria offer no additional benefit in this regime. The practical value of the dual criterion emerges only at the neural-network level, where approximate κ estimation introduces false positives (cf. Conjecture 3).

D Negative Result: Gram Point Curvature

A natural question is whether the curvature κ at Gram points g_n (defined by $\vartheta(g_n) = n\pi$) carries a signature of Gram's law violations. We computed $\kappa(g_n + 0.03)$ for $n = 0, \ldots, 200$ and classified each Gram point as "good" (satisfying Gram's law) or "bad" (violating it). Only 3 out of 201 Gram points in this range are bad ($n \in \{126, 134, 195\}$), yielding insufficient statistical power for any rank-based test (Mann–Whitney requires $n \geq 8$ per group). This experiment is therefore **negative** — not because the hypothesis is disproved, but because the data are inadequate to test it. Extending to $n > 1000$, where bad Gram points become more frequent, would be needed to revisit this question. [negative, 2026-04-14]

E Notation Index

Symbol	Domain	Meaning
R_q	QROP-Net	Polynomial ring $\mathbb{Z}_q[X]/(X^n+1)$
ϕ_i	QROP-Net	Optical phase of ciphertext element i
U_0	QROP-Net	Transmitted complex wavefront $e^{i\phi}$
d	QROP-Net	Learnable propagation distance (mm)
α	QROP-Net	FrFT order angle (radians)
$\xi(s)$	Resonance	Completed Riemann zeta function
$\Phi_\xi(t)$	Resonance	Phase of $\xi(1/2+it)$
$A(t)$	Resonance	Amplitude $\lvert\xi(1/2+it)\rvert$
$\Psi(n)$	Resonance	Smooth prime-counting phase
σ^*	Resonance	Real part of resonant point
ℓ	Resonance	σ-lift: $\sigma^*-1/2$
$F_2(t)$	Polyadic	Final discriminant function
$\hat{L}(t)$	Polyadic	Central-difference log-derivative
$v(\sigma,t)$	Both	Phase of $\xi(\sigma+it)$ / general phase
$\mathcal{P}(\phi;L)$	Both	Phase quantisation operator $\sin^2(L\phi/2)$